

# Coping with the overfitting issue: regularizing and pruning irrelevant hidden nodes

國立政治大學 資訊管理學系

蔡瑞煌 特聘教授

# AI applications

- Training phase: (training) data + AI model + algorithm & code + setting of network & hyperparameters → AI model/AI system
- Inferencing phase: performance is obtained from model((test) data)
- Goals of training are reasonable inferencing

Where we are now...

ideas/concepts →  
modules →  
learning algorithm →  
codes →  
intelligent systems

Where we are now...

# Types of Learning

**Supervised:** Learning with a **labeled training set**


Example: email *classification* with already labeled emails

**Unsupervised:** Discover **patterns** in **unlabeled data**

Example: *cluster* similar documents based on text

**Reinforcement learning:** learn to **act** based on **feedback/reward**


Example: learn to *play* Go, reward: *win or lose*




Classification

Anomaly Detection  
Sequence labeling

...



Regression



Clustering

<http://mbjoseph.github.io/2013/11/27/measure.html>

Where we are now...

# The supervised learning problems: Regression and Classification




## Regression

What is the temperature going to be tomorrow?



## Classification

Will it be Cold or Hot tomorrow?



Where we are now...

# Stopping criteria (also the learning goals) for **regression** applications

one output node

The learning process should stop when

~~1.  $L_N(\mathbf{w}) = 0$~~

2. a tiny  $L_N(\mathbf{w})$  value

3.  $|f(\mathbf{x}^c, \mathbf{w}) - y^c| < \varepsilon \quad \forall c$  with  $\varepsilon$  being tiny


$$L_N(\mathbf{w}) \equiv \frac{1}{N} \sum_{c=1}^N (f(\mathbf{x}^c, \mathbf{w}) - y^c)^2$$

- Each reasonable learning goal can be used as a stopping criterion.
- Different stopping criterion results in different length of training time and different model.


# The regression applications

## The learning goal

$$|f(\mathbf{x}^c, \mathbf{w}) - y^c| \leq \varepsilon \quad \forall \quad c \in \mathbf{I}$$



## The inferencing mechanism




Where we are now...

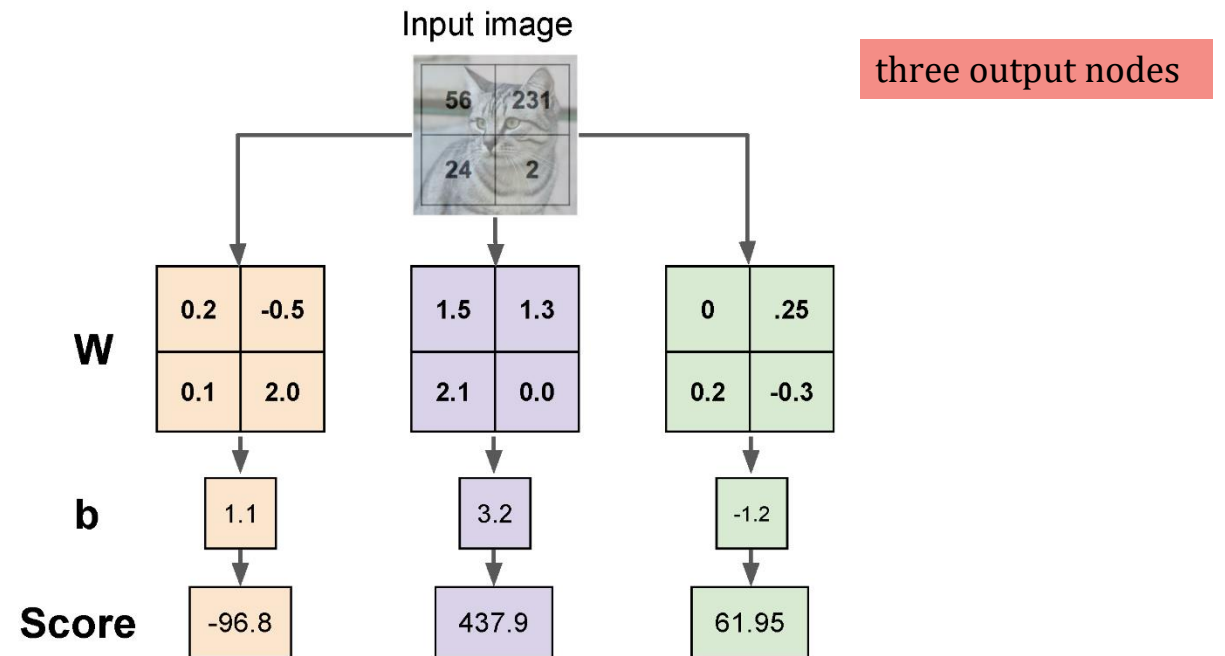
# The **three-class classification** applications

Example with an image with 4 pixels, and 3 classes (cat/dog/ship)

## Algebraic Viewpoint



## Visual Viewpoint

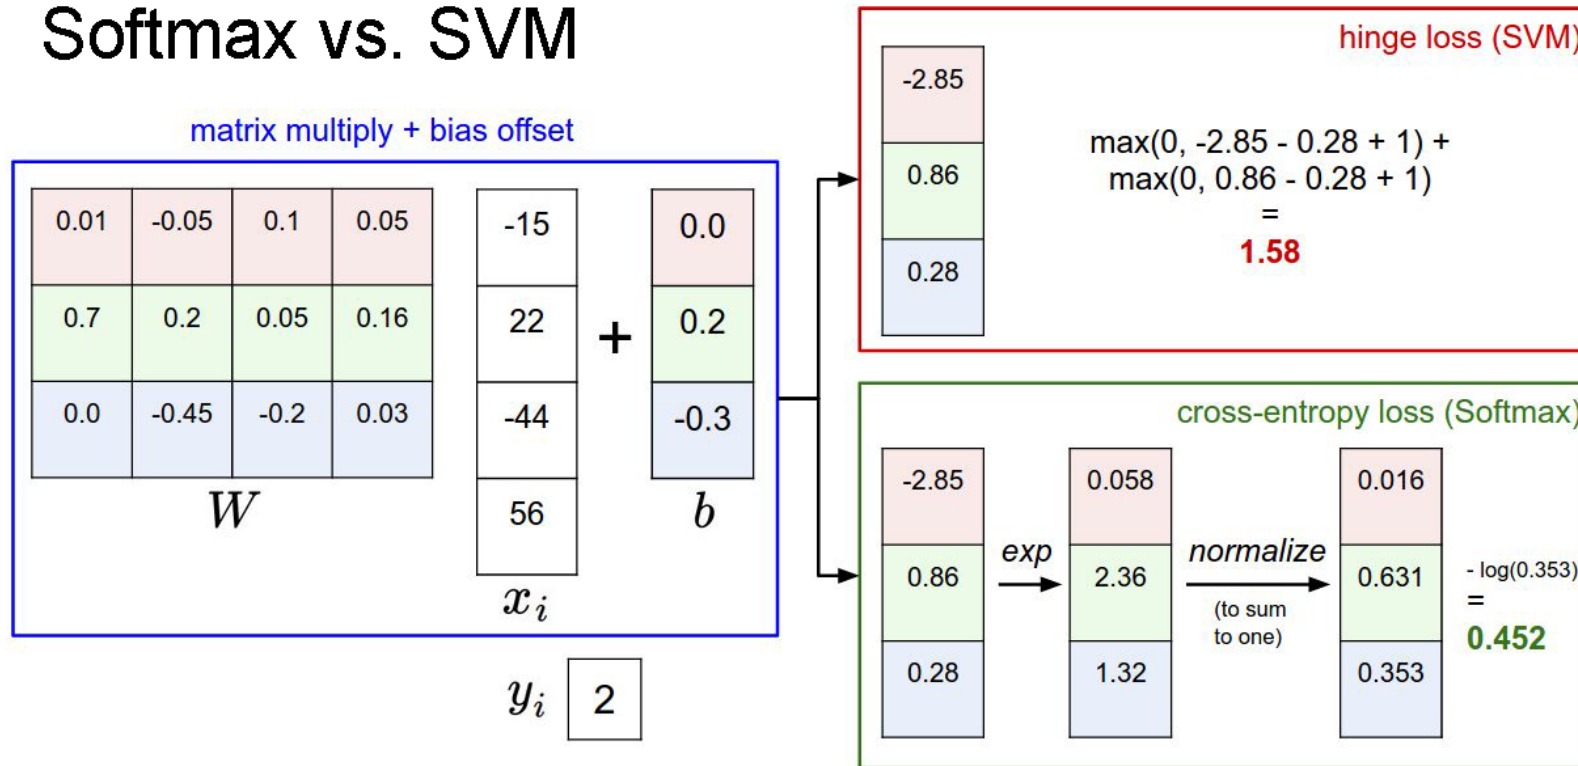




Where we are now...

# The **three-class classification** applications

## Softmax vs. SVM



# Classification Applications

## Design (x attributes & y label)

**x:**

- ✓ 年齡
- ✓ 性別
- ✓ 診斷
- ✓ 分期
- ✓ 診斷癌症期間
- ✓ 復發與否
- ✓ Kps (身體功能)
- ✓ gs1-gs22 (22項)症狀 (0:無; 1:有)

**y:** 疲倦，個案總計 686位，過去一周平均疲倦程度 (f3)  
(0-10分，real-value variable) → 過去一周疲倦分組  
(Gf3) (分成三組，binary variable) :

- ✓ 無: 46位
- ✓ 輕度: 346 位
- ✓ 中至重度: 294位

# Classification Applications Design ( $y$ label)

Output value: real number

*SLFN with one output node and linear (output) function*

疲倦

- ✓ 無: 46位
- ✓ 輕度: 346位
- ✓ 中至重度: 294位

Learning phase:

$y$  (i.e., target output):

- ✓ 無: 0
- ✓ 輕度: 5
- ✓ 中至重度: 10

Inferencing phase:

$f$  (i.e., actual output):

- ✓  $[-2.5, 2.5) \rightarrow$  無
- ✓  $[2.5, 7.5) \rightarrow$  輕度
- ✓  $[7.5, 12.5) \rightarrow$  中至重度
- ✓  $(-\infty, -2.5) \text{ OR } [12.5, \infty) \rightarrow$  unknown

Output value: binary number

*SLFN with three output nodes and softmax arrangement*

疲倦

- ✓ 無: 46位
- ✓ 輕度: 346位
- ✓ 中至重度: 294位

Learning phase:

$y$  (i.e., target output):

- ✓ 無: (1, 0, 0)
- ✓ 輕度: (0, 1, 0)
- ✓ 中至重度: (0, 0, 1)

Inferencing phase:

$f$  (i.e., actual output):

- ✓ (1, 0, 0)  $\rightarrow$  無
- ✓ (0, 1, 0)  $\rightarrow$  輕度
- ✓ (0, 0, 1)  $\rightarrow$  中至重度

# Another classification application

## Design (x attributes & y label)

- Data

x attributes	
x1	性別
x2	年齡
x3	國籍
x4	婚姻狀態
x5	直系親屬數
x6	最高學歷
x7	來台時長
x8	平均月收入
x9	剩餘居留時間

x attributes	
x10	借款時長
x11	借款金額
x12	用途
x13	工作性質
x14	工作地點
x15	雇主資訊
x16	薪資如期撥入
x17	薪資撥付方式
x18	薪資結匯方式

y label	
y (target output, real number)	信用評級 有5個等級

y	信用評級
1	E最差
2	D
3	C
4	B
5	A最好

# Another classification application

## Design ( $y$ label)

- $y$  (target output)  $\in \{1, 2, 3, 4, 5\}$  one output node
- At the learning phase, let  $\varepsilon = 0.2$ . Then the learning goal is to make  $f$  (actual output)  $\in \{[0.8, 1.2], [1.8, 2.2], [2.8, 3.2], [3.8, 4.2], [4.8, 5.2]\}$ .
- At the inferencing phase,  $y = 1$  if  $f \in [0.5, 1.5)$ ;  $y = 2$  if  $f \in [1.5, 2.5)$ ;  $y = 3$  if  $f \in [2.5, 3.5)$ ;  $y = 4$  if  $f \in [3.5, 4.5)$ ;  $y = 5$  if  $f \in [4.5, 5.5)$
- $y$  is unknown if  $f < 0.5$  OR  $f \geq 5.5$ .

Where we are now...

Stopping criteria (also the learning goals) for the SLFN with  
**each output node** whose output values are **real numbers**  
for the **two-class classification** application


- **Two-class classification** problems with  $\mathbf{I} \equiv \mathbf{I}_1 \cup \mathbf{I}_2$ , where  $\mathbf{I}_1$  and  $\mathbf{I}_2$  are the sets of indices of given cases in **classes 1 and 2**. Furthermore,  $y^c$  is the target of the  $c^{\text{th}}$  case, with **1 and 0** being the targets of classes 1 and 2
- When the SLFN with **only one output node** whose output value is **real number**, the stopping criteria may be as follows:
  1.  $|f(\mathbf{x}^c, \mathbf{w}) - y^c| < \varepsilon \quad \forall c$
  2.  $f(\mathbf{x}^c, \mathbf{w}) > \nu \quad \forall c \in \mathbf{I}_1$  and  $f(\mathbf{x}^c, \mathbf{w}) \leq -\nu \quad \forall c \in \mathbf{I}_2$ , with  $1 > \nu > 0$
  3.  $\alpha \equiv \min_{c \in \mathbf{I}_1} f(\mathbf{x}^c, \mathbf{w}) > \beta \equiv \max_{c \in \mathbf{I}_2} f(\mathbf{x}^c, \mathbf{w})$   
(Linearly separating condition, *LSC*)

Different stopping criterion results  
in different length of training time  
and different model.

Where we are now...


# Stopping criteria (also the learning goals) for the SLFN with each output node whose output values are real numbers for the two-class classification application

$$y^c = 1 \ \forall \ c \in \mathbf{I}_1; y^c = 0 \ \forall \ c \in \mathbf{I}_2 \quad \text{X} : f(\mathbf{x}^c, \mathbf{w}), \ \forall \ c \in \mathbf{I}_1 \quad \text{O} : f(\mathbf{x}^c, \mathbf{w}), \ \forall \ c \in \mathbf{I}_2$$




learning goal type 1  
(also inferencing goal):  
 $|f(\mathbf{x}^c, \mathbf{w}) - y^c| \leq \varepsilon \ \forall \ c \in \mathbf{I}_1;$   
 $|f(\mathbf{x}^c, \mathbf{w}) + y^c| \leq \varepsilon \ \forall \ c \in \mathbf{I}_2$

$\varepsilon$  is a hyperparameter regarding the learning!



learning goal type 2  
(also inferencing goal):  
 $f(\mathbf{x}^c, \mathbf{w}) \geq v \ \forall \ c \in \mathbf{I}_1;$   
 $f(\mathbf{x}^c, \mathbf{w}) \leq -v \ \forall \ c \in \mathbf{I}_2$

$v$  is a hyperparameter regarding the inferencing!




learning goal type 3: LSC

Where we are now...

Stopping criteria (also the learning goals) for the SLFN with  
**each output node** whose output values are **real numbers**  
for the **two-class classification** application


$$y^c = 1 \ \forall c \in \mathbf{I}_1; y^c = 0 \ \forall c \in \mathbf{I}_2 \quad \text{X} : f(\mathbf{x}^c, \mathbf{w}), \ \forall c \in \mathbf{I}_1 \quad \text{O} : f(\mathbf{x}^c, \mathbf{w}), \ \forall c \in \mathbf{I}_2$$

- $|f(\mathbf{x}^c, \mathbf{w}) - y^c| < \varepsilon \ \forall c$



learning goal type 1  
(also inferencing goal):  
 $|f(\mathbf{x}^c, \mathbf{w}) - 1| \leq \varepsilon \ \forall c \in \mathbf{I}_1;$   
 $|f(\mathbf{x}^c, \mathbf{w})| \leq \varepsilon \ \forall c \in \mathbf{I}_2$

$\varepsilon$  is a hyperparameter



The inferencing  
mechanism




Where we are now...

Stopping criteria (also the learning goals) for the SLFN with  
**each output node** whose output values are **real numbers**  
for the **two-class classification** application

$$y^c = 1 \ \forall c \in \mathbf{I}_1; y^c = 0 \ \forall c \in \mathbf{I}_2 \quad \textcolor{blue}{\mathbf{x}} : f(\mathbf{x}^c, \mathbf{w}), \ \forall c \in \mathbf{I}_1 \quad \textcolor{red}{\mathbf{o}} : f(\mathbf{x}^c, \mathbf{w}), \ \forall c \in \mathbf{I}_2$$

- $f(\mathbf{x}^c, \mathbf{w}) \geq \nu \ \forall c \in \mathbf{I}_1$  and  $f(\mathbf{x}^c, \mathbf{w}) \leq -\nu \ \forall c \in \mathbf{I}_2$  with  $1 > \nu > 0$ .




Where we are now...

Stopping criteria (also the learning goals) for the SLFN with  
**each output node** whose output values are **real numbers**  
for the **two-class classification** application


$$y^c = 1 \ \forall c \in \mathbf{I}_1; y^c = 0 \ \forall c \in \mathbf{I}_2 \quad \text{X} : f(\mathbf{x}^c, \mathbf{w}), \ \forall c \in \mathbf{I}_1 \quad \text{O} : f(\mathbf{x}^c, \mathbf{w}), \ \forall c \in \mathbf{I}_2$$

- The LSC (Tsaih, 1993)



$$\alpha \equiv \min_{c \in \mathbf{I}_1} f(\mathbf{x}^c, \mathbf{w}); \beta \equiv \max_{c \in \mathbf{I}_2} f(\mathbf{x}^c, \mathbf{w})$$

learning goal type 3: LSC




The inferencing mechanism:

$$f(\mathbf{x}^c, \mathbf{w}) \geq v \ \forall c \in \mathbf{I}_1; \\ f(\mathbf{x}^c, \mathbf{w}) \leq -v \ \forall c \in \mathbf{I}_2$$

Where we are now...

Stopping criteria (also the learning goals) for the SLFN with  
**each output node** whose output values are **real numbers**  
for the **two-class classification** application



$$\alpha \equiv \min_{c \in \mathbf{I}_1} f(\mathbf{x}^c, \mathbf{w}); \beta \equiv \max_{c \in \mathbf{I}_2} f(\mathbf{x}^c, \mathbf{w})$$

learning goal type 3: LSC

When LSC ( $\alpha > \beta$ ) is true, the inferencing mechanism

$$f(\mathbf{x}^c, \mathbf{w}) \geq v \quad \forall c \in \mathbf{I}_1 \text{ and } f(\mathbf{x}^c, \mathbf{w}) \leq -v \quad \forall c \in \mathbf{I}_2$$

can be set by directly adjusting  $\mathbf{w}^o$  according to the following formula:

$$\frac{2v}{\alpha - \beta} w_i^o \rightarrow w_i^o \quad \forall i,$$


The weight vector between the hidden layer and the output node

$$\text{then } v - \min_{c \in \mathbf{I}_1} \sum_{i=1}^p w_i^o a_i^c \rightarrow w_0^o$$

The threshold of the output node

Where we are now...

# Learning dilemma of gradient-descent-based learning



# Extra stopping criteria for the learning (not the learning goals)

$\|\nabla_{\mathbf{w}} L_N(\mathbf{w})\|$  is the length of  $\nabla_{\mathbf{w}} L_N(\mathbf{w})$ .

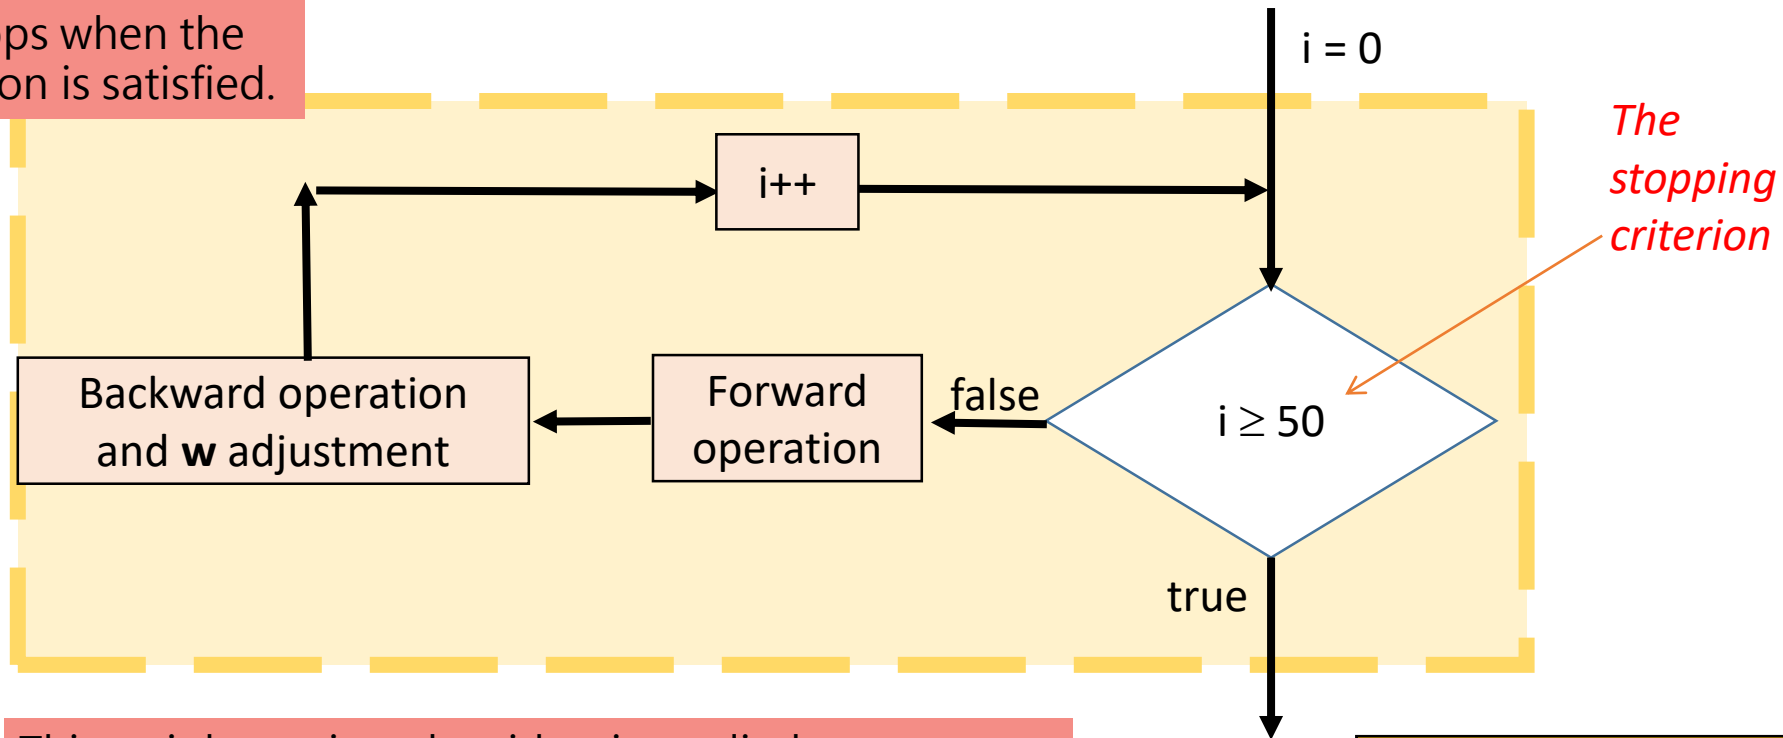
- ~~1. The learning process should stop when  $\|\nabla_{\mathbf{w}} L_N(\mathbf{w})\| = 0$  but a tiny  $L_N(\mathbf{w})$  value cannot be accomplished.~~
2. The learning process should stop when  $\|\nabla_{\mathbf{w}} L_N(\mathbf{w})\|$  is tiny but a tiny  $L_N(\mathbf{w})$  value cannot be accomplished.
3. The learning process should stop when  $\eta$  (the learning rate) is tiny but a tiny  $L_N(\mathbf{w})$  value cannot be accomplished.

The neighborhood of undesired attractors, where  $\|\nabla_{\mathbf{w}} L_N(\mathbf{w})\| \approx 0$  but a tiny  $L_N(\mathbf{w})$  value cannot be accomplished:

- a) the local minimum, the saddle point, or the plateau.
- b) the global minimum of the defective network architecture.

# The flowchart of **weight-tuning** algorithms for 2-layer neural networks in CS231n

The process stops when the stopping criterion is satisfied.



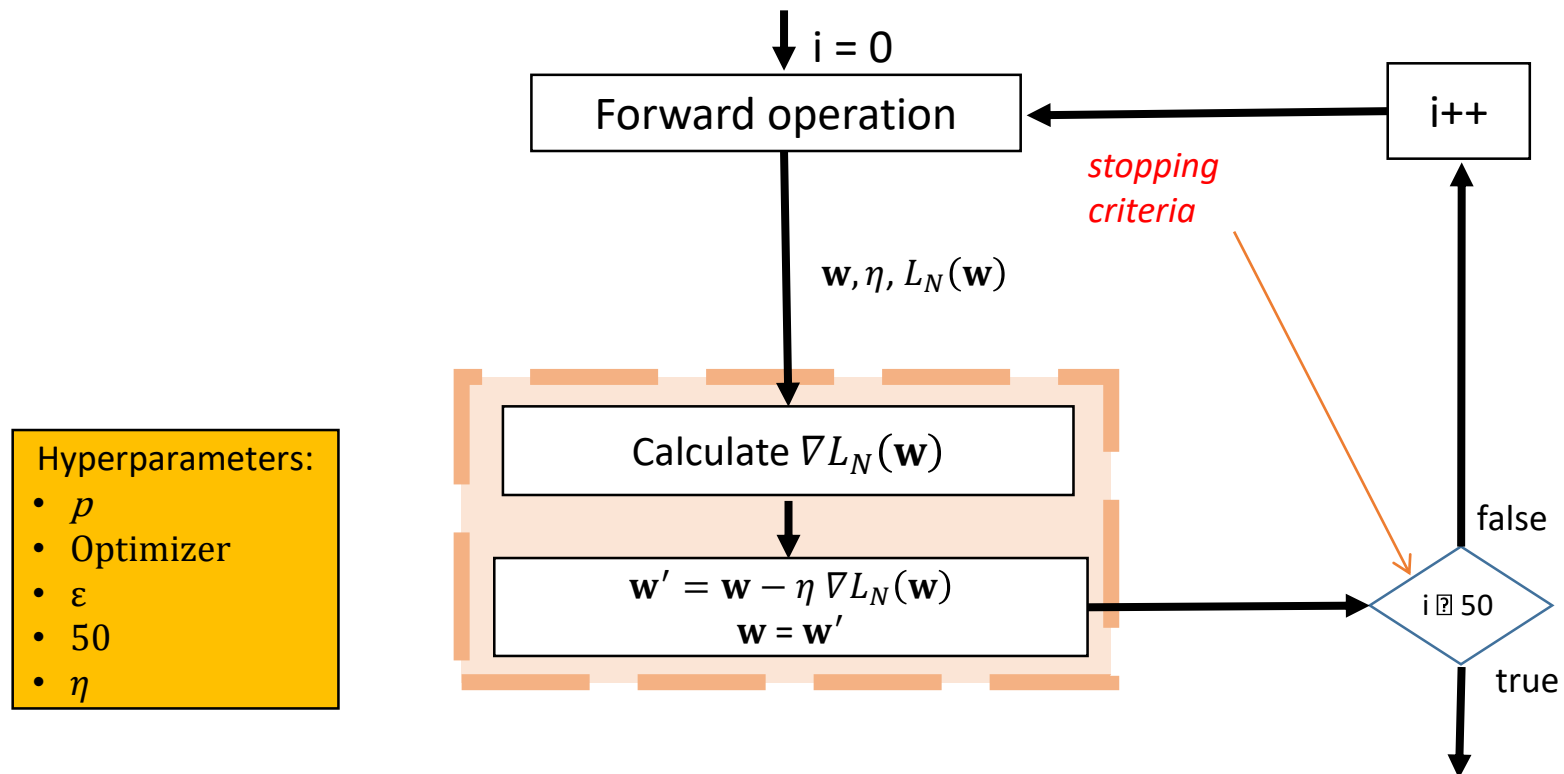
This weight-tuning algorithm is applied to many kinds of Neural Networks, including 2-layer neural networks, CNN, RNN, reinforcement learning, GAN, BERT, and so on.

## Hyperparameters:

- $p$
- Optimizer: SGD
- Epoch upper bound: 50
- Learning rate:  $1e-6$

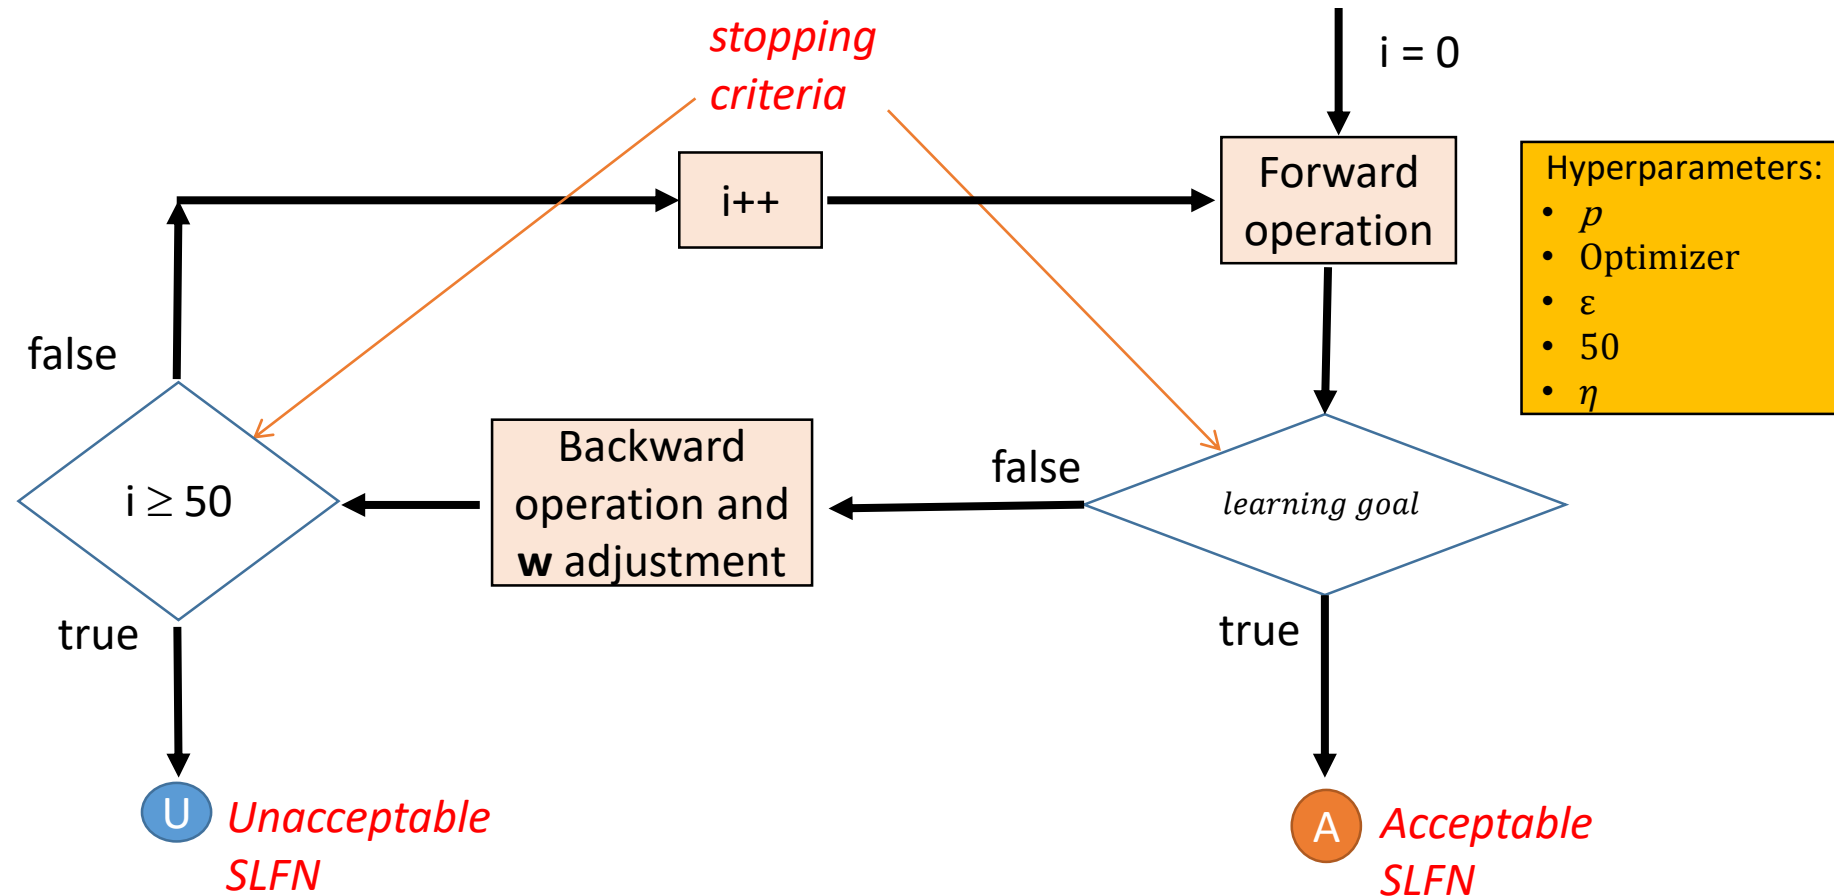
Where we are now...

# The flowchart of **weight-tuning** module\_EU



Where we are now...

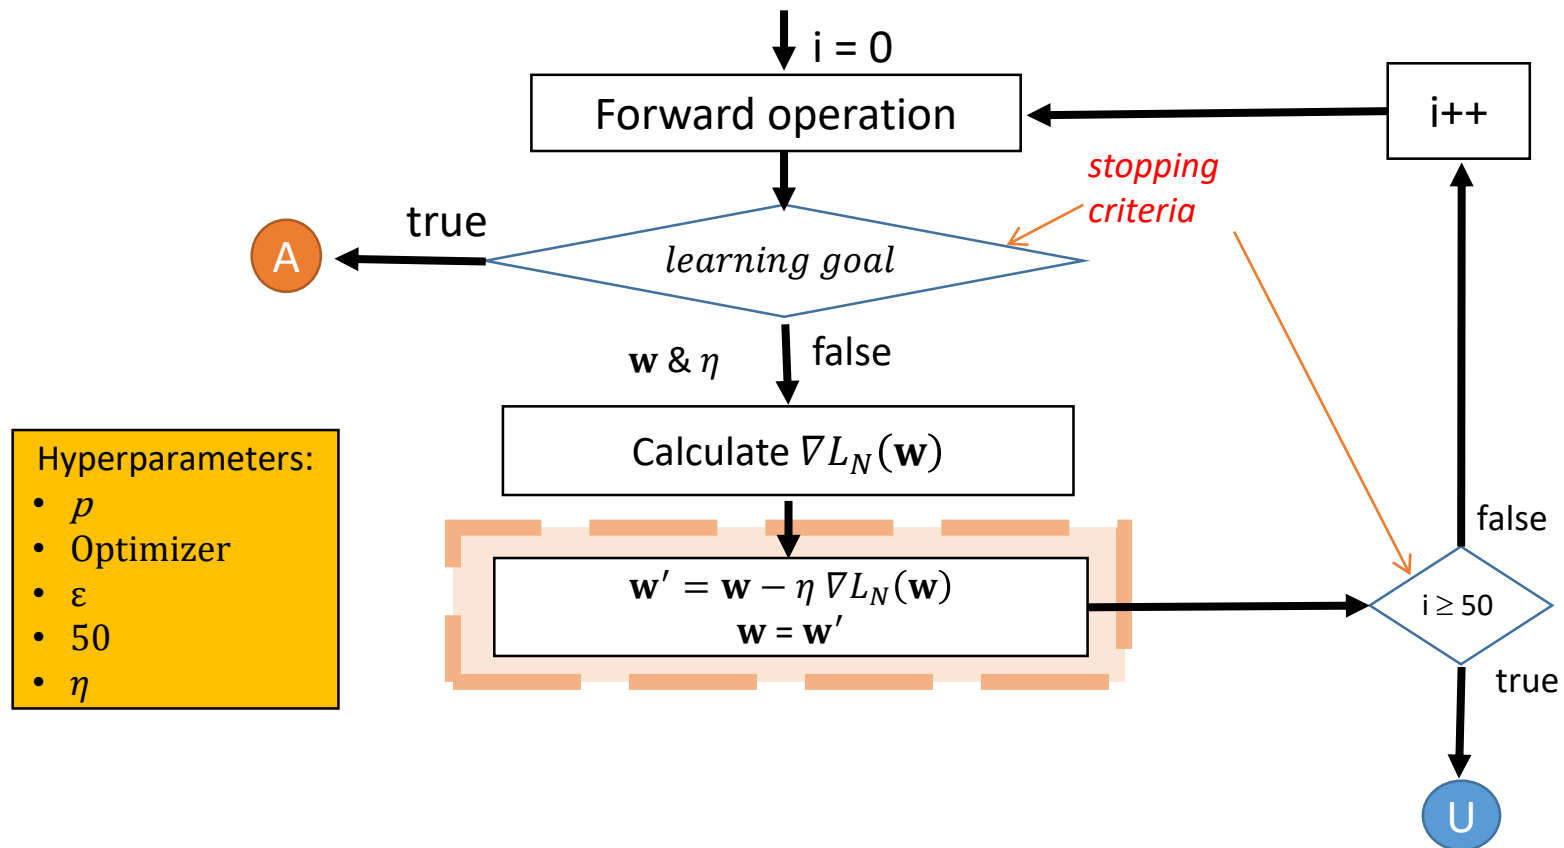
The flowchart of **weight-tuning** module\_EU\_LG that indicate either an unacceptable SLFN or an acceptable SLFN





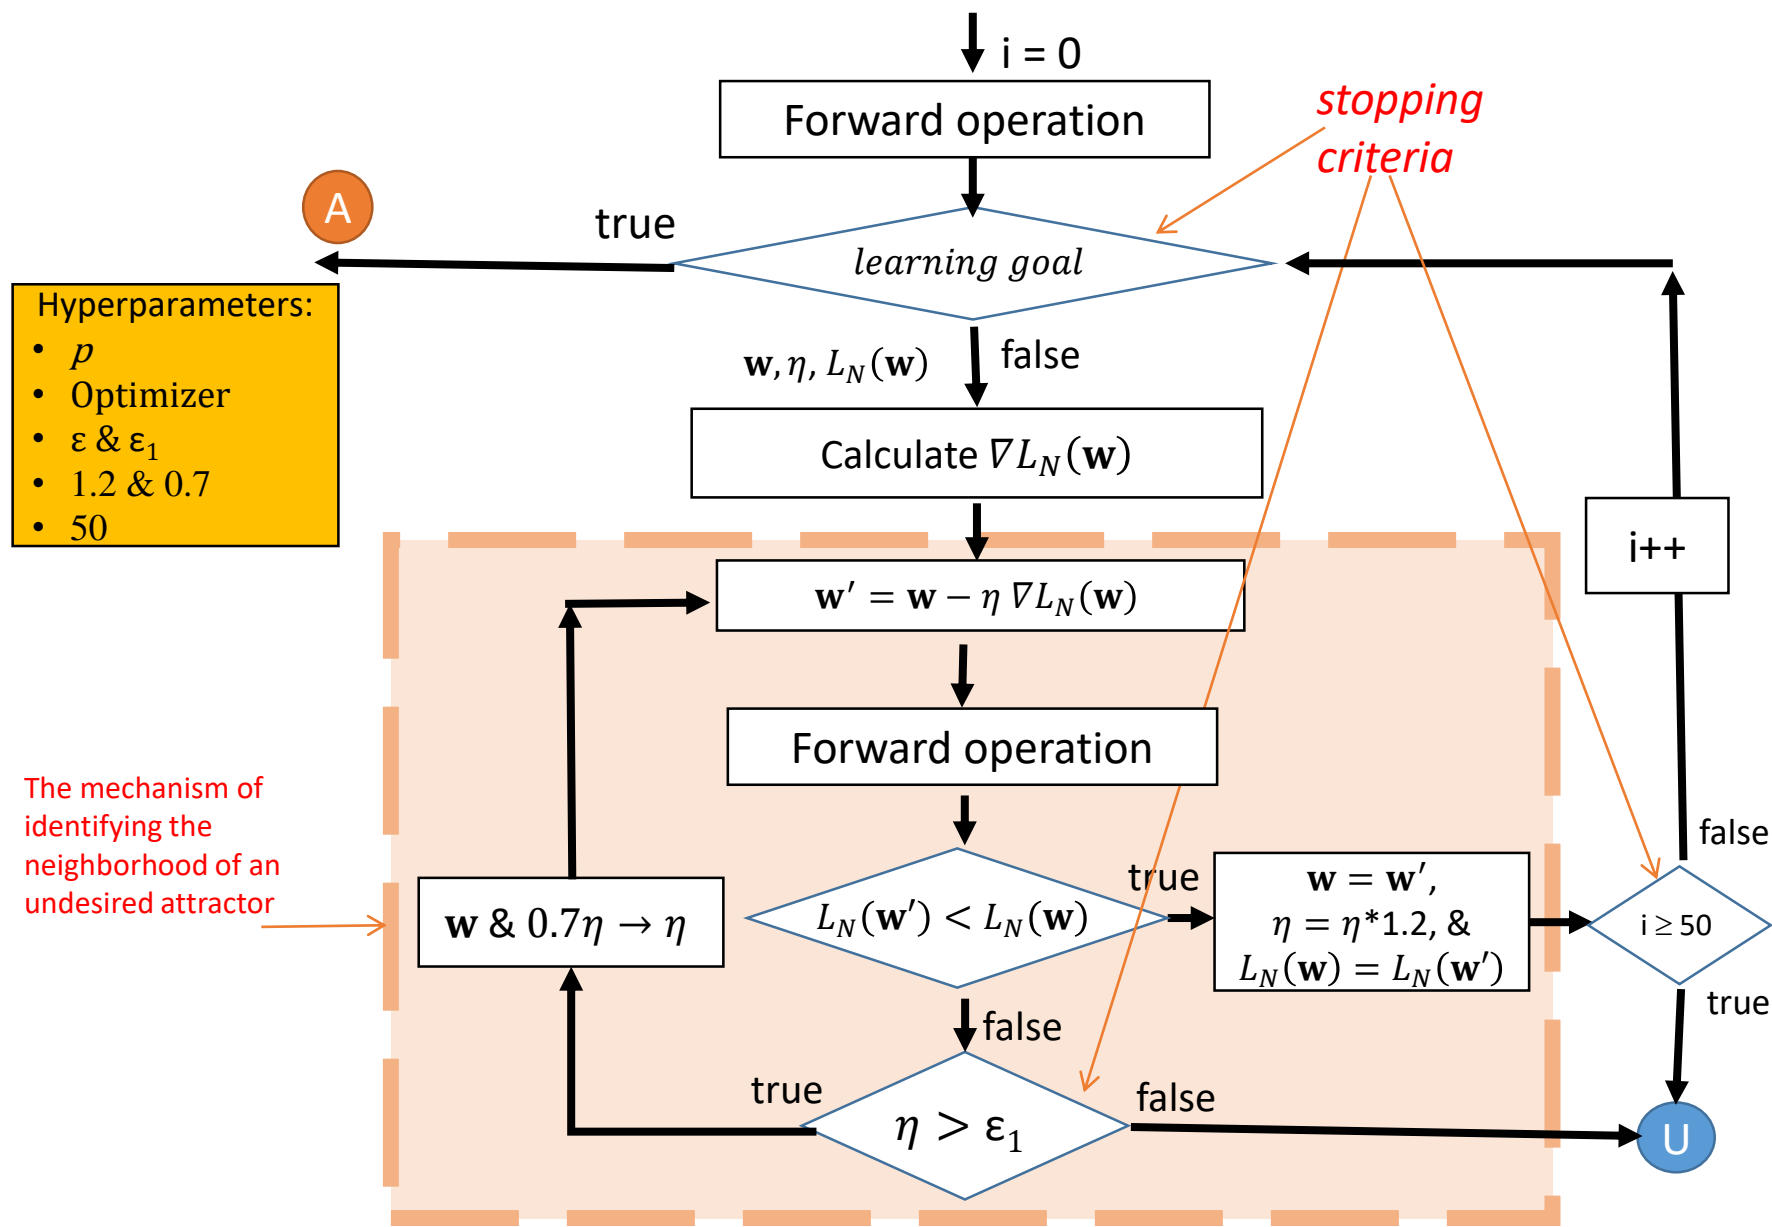
Where we are now...

# The flowchart of **weight-tuning** module\_EU\_LG



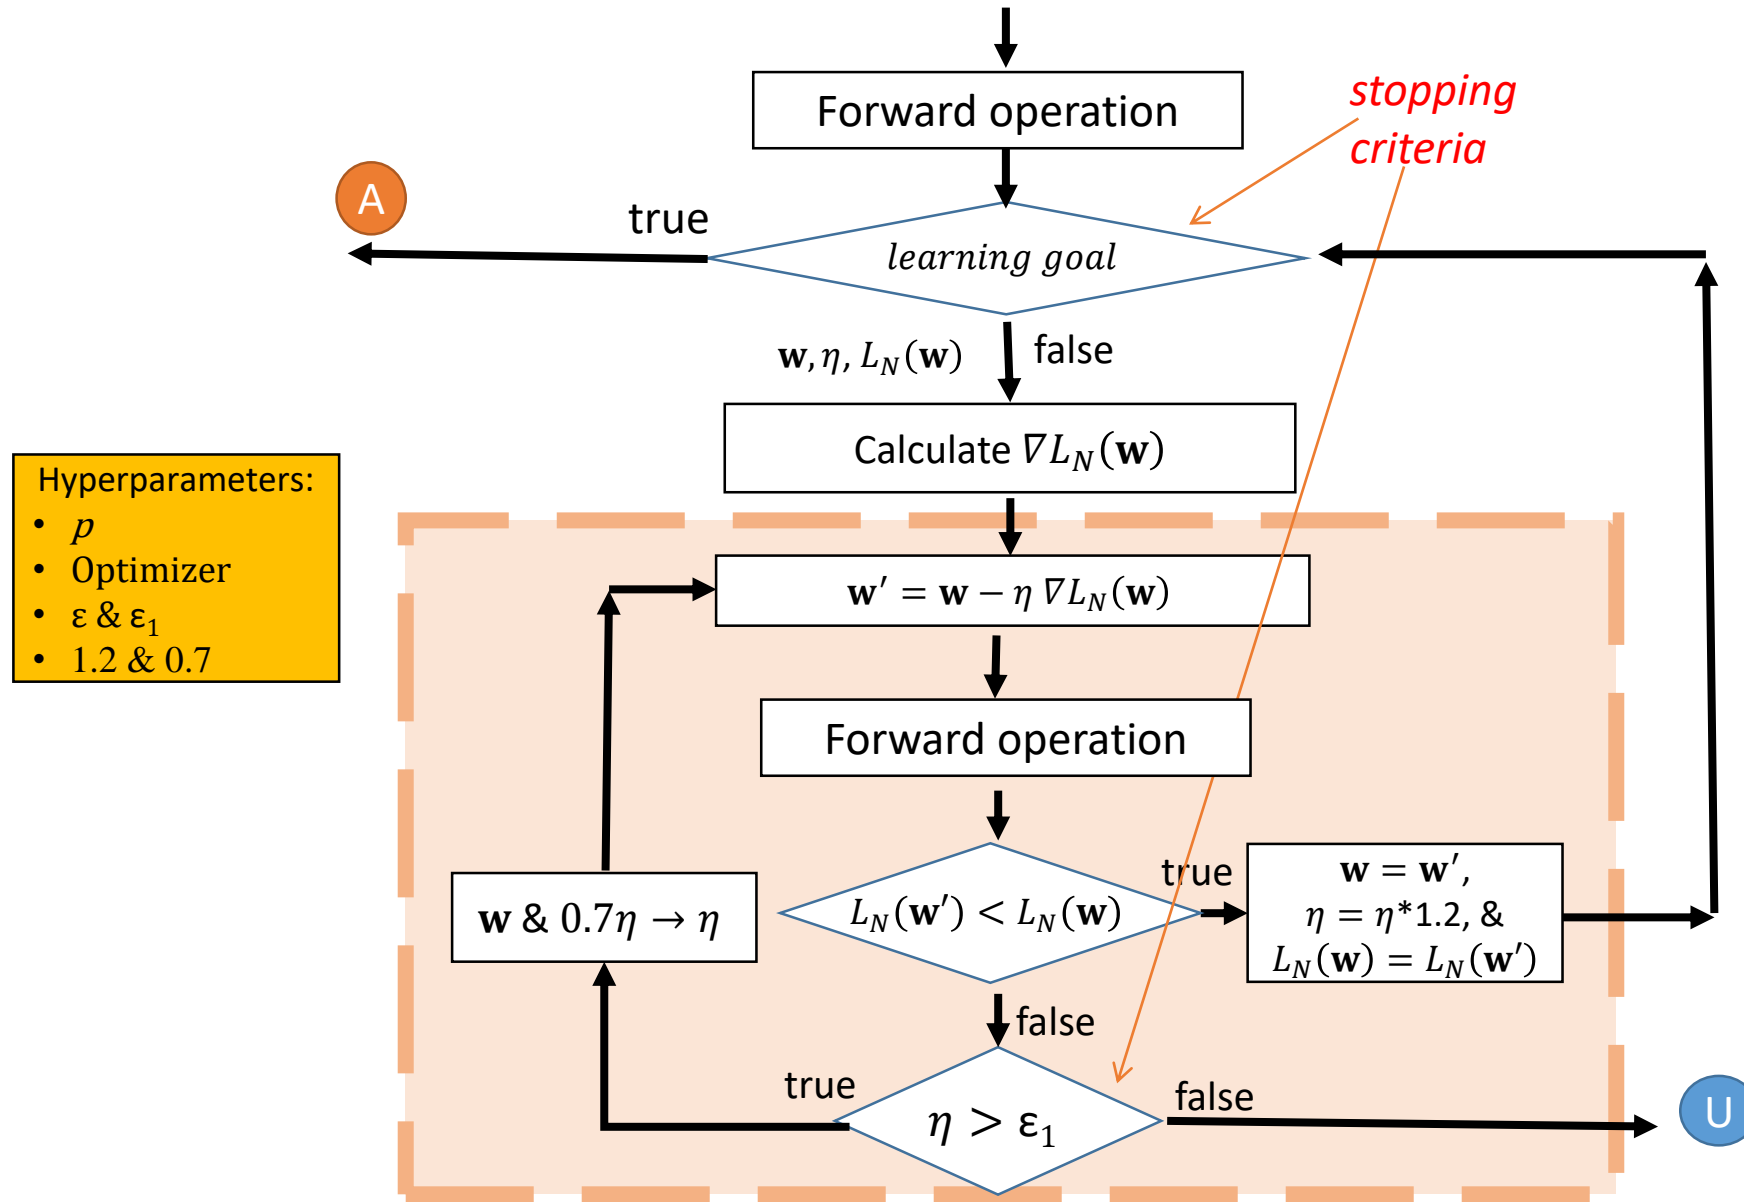
Where we are now...

# The flowchart of **weight-tuning** module\_EU\_LG\_UA



Where we are now...

# The flowchart of **weight-tuning** module\_LG\_UA



# Performance differences amongst weight-tuning modules

- There are four weight-tuning modules
  - ✓ the weight-tuning module\_EU  
The simplest and the learning time length is expected
  - ✓ the weight-tuning module\_EU\_LG  
Shorter learning time length than the weight-tuning module\_EU
  - ✓ the weight-tuning module\_EU\_LG\_UA  
The learning time length may be longer than the weight-tuning module\_EU\_LG
  - ✓ the weight-tuning module\_LG\_UA  
The learning time length is not an issue

# Algorithm representation and development

([Algorithm - Wikipedia](#))


- Algorithms can be expressed in many kinds of notation, including **natural languages**, **pseudocode**, **flowcharts**, drakon-charts, **programming languages** or control tables (processed by interpreters).
  - ✓ Natural language expressions of algorithms tend to be **verbose and ambiguous**, and are rarely used for complex or technical algorithms.
  - ✓ Pseudocode, flowcharts, drakon-charts and control tables are **structured ways** to express algorithms that avoid many of the ambiguities common in the statements based on natural language.
  - ✓ Programming languages are primarily intended for expressing algorithms in **a form that can be executed by a computer**, but are also often used as a way to define or document algorithms.
- Typical steps in the development of algorithms:
  - ✓ Problem definition ← **learning-based prediction problem**
  - ✓ Development of a model ← **2-layer net, 4-layer net, or deep neural networks**
  - ✓ Specification of the algorithm ← **The learning algorithm**
  - ✓ Designing an algorithm ← **The gradient-descent-based learning algorithm**
  - ✓ Checking the correctness of the algorithm ← **The math proof of the proposed learning algorithm**
  - ✓ Analysis of algorithm ← **The amount of parameters, the (learning and inferencing) time scale, ...**
  - ✓ Implementation of algorithm ← **The coding**
  - ✓ **Program testing**
  - ✓ Documentation preparation

# Program testing -- Performance of AI Applications


- How do AI professionals evaluate the performance of the AI applications?  
    ← effectiveness & efficiency
- However, there are learning dilemma and overfitting when evaluating the effectiveness & efficiency.
- You need to deal with learning dilemma and overfitting, not only for the purposes of learning, but also of inferencing.

Where we are now...


# overfitting




$\rightarrow \theta_0 + \theta_1 x$   
"Underfit" "High bias"



$\rightarrow \theta_0 + \theta_1 x + \theta_2 x^2$   
"Just right"



$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$   
"Overfit" "High variance"



inadequate

good compromise

over-fitting

# Generalization

- Learned hypothesis may fit the training data very well, even noises (or outliers) in the training data, but fail to generalize to new examples (test data)
- In machine learning and statistical regression, the generalization error (also known as the out-of-sample error) is a measure of how accurately an algorithm is able to predict outcome values for previously unseen data.

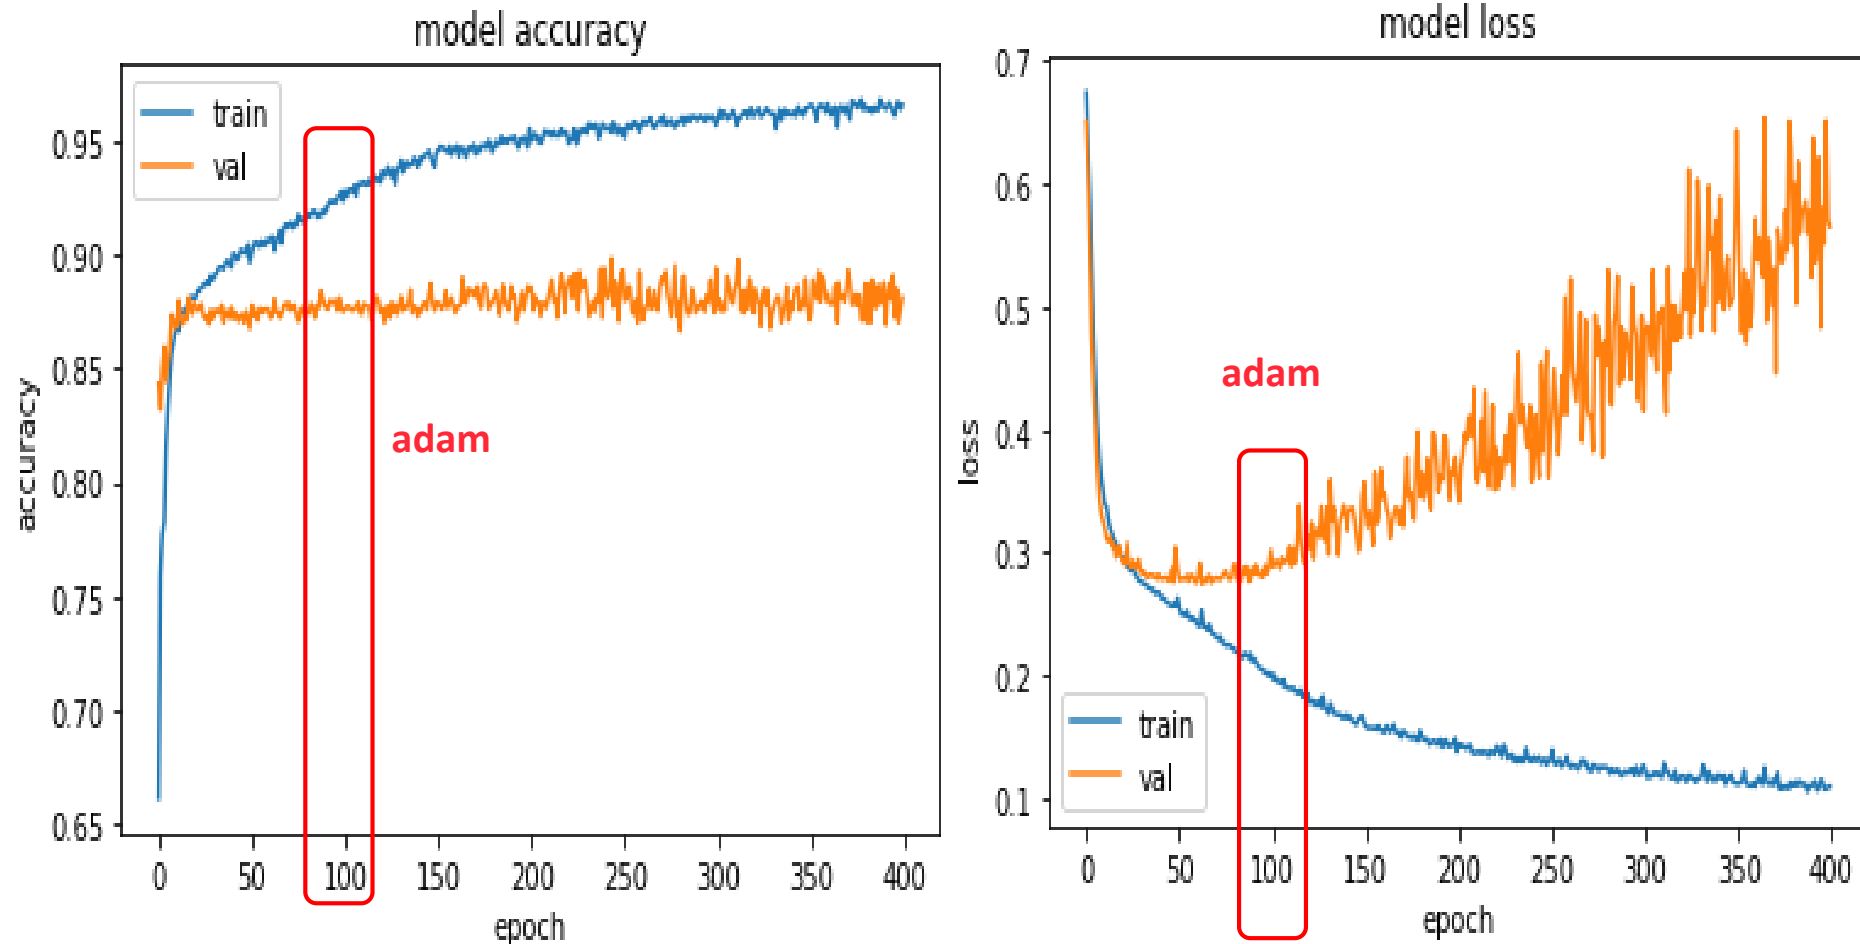


# Learning curves

- Because learning algorithms are evaluated on **finite samples**, the evaluation of a learning algorithm may be sensitive to **sampling error**.
- As a result, measurements of prediction error on the current data may not provide much information about predictive ability on new data.
- The performance of a learning algorithm is measured by **plots of the generalization error values** through the learning process, which are called **learning curves**.
- Generalization error can be minimized by avoiding **overfitting** in the learning process.

Where we are now...

# Learning curve and overfitting

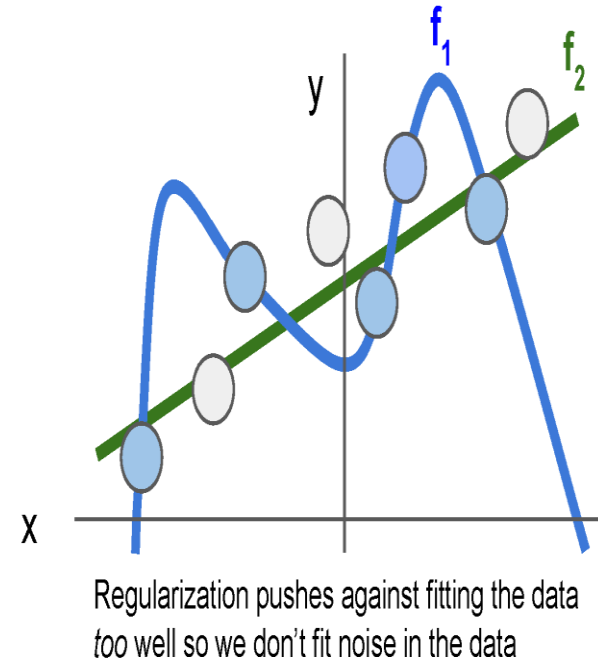


Early stop!

# Overfitting

In **statistics**, **overfitting** is "the production of an analysis that corresponds **too closely or exactly** to a particular set of data, and may therefore **fail to fit additional data** or predict future observations **reliably**."


Regularization: Prefer Simpler Models



# Overfitting

An **over-fitted model** is a model that contains **more parameters** than can be justified by the data.

Regularization: Prefer Simpler Models



Regularization pushes against fitting the data too well so we don't fit noise in the data

# Regularization

$\lambda$  = regularization strength  
(hyperparameter)

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Data loss: Model predictions should match training data}} + \underbrace{\lambda R(W)}_{\text{Regularization: Prevent the model from doing too well on training data}}$$

**Data loss:** Model predictions should match training data

**Regularization:** Prevent the model from doing *too* well on training data

## Simple examples

L2 regularization:  $R(W) = \sum_k \sum_l W_{k,l}^2$

L1 regularization:  $R(W) = \sum_k \sum_l |W_{k,l}|$

Elastic net (L1 + L2):  $R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$

## More complex:

Dropout

Batch normalization

Stochastic depth, fractional pooling, etc

# Regularization

$\lambda$  = regularization strength  
(hyperparameter)

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Data loss}} + \underbrace{\lambda R(W)}_{\text{Regularization}}$$

**Data loss:** Model predictions should match training data

**Regularization:** Prevent the model from doing *too* well on training data

Why regularize?

- Express preferences over weights
- Make the model *simple* so it works on test data
- Improve optimization by adding curvature

# Regularization - In practice

**Training:** Add random noise

**Testing:** Marginalize over the noise

## **Examples:**

Dropout

Batch Normalization

Data Augmentation

DropConnect

Fractional Max Pooling

Stochastic Depth

Cutout / Random Crop

Mixup

- Consider dropout for large fully-connected layers
- Batch normalization and data augmentation almost always a good idea
- Try cutout and mixup especially for small classification datasets

# Summary: the overfitting may be due to **big weights**

## Regularization

$\lambda$  = regularization strength  
(hyperparameter)

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Data loss: Model predictions should match training data}} + \underbrace{\lambda R(W)}_{\text{Regularization: Prevent the model from doing too well on training data}}$$

**Data loss:** Model predictions should match training data

**Regularization:** Prevent the model from doing *too* well on training data

### Simple examples

L2 regularization:  $R(W) = \sum_k \sum_l W_{k,l}^2$

L1 regularization:  $R(W) = \sum_k \sum_l |W_{k,l}|$

Elastic net (L1 + L2):  $R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$

### More complex:


Dropout

Batch normalization

Stochastic depth, fractional pooling, etc

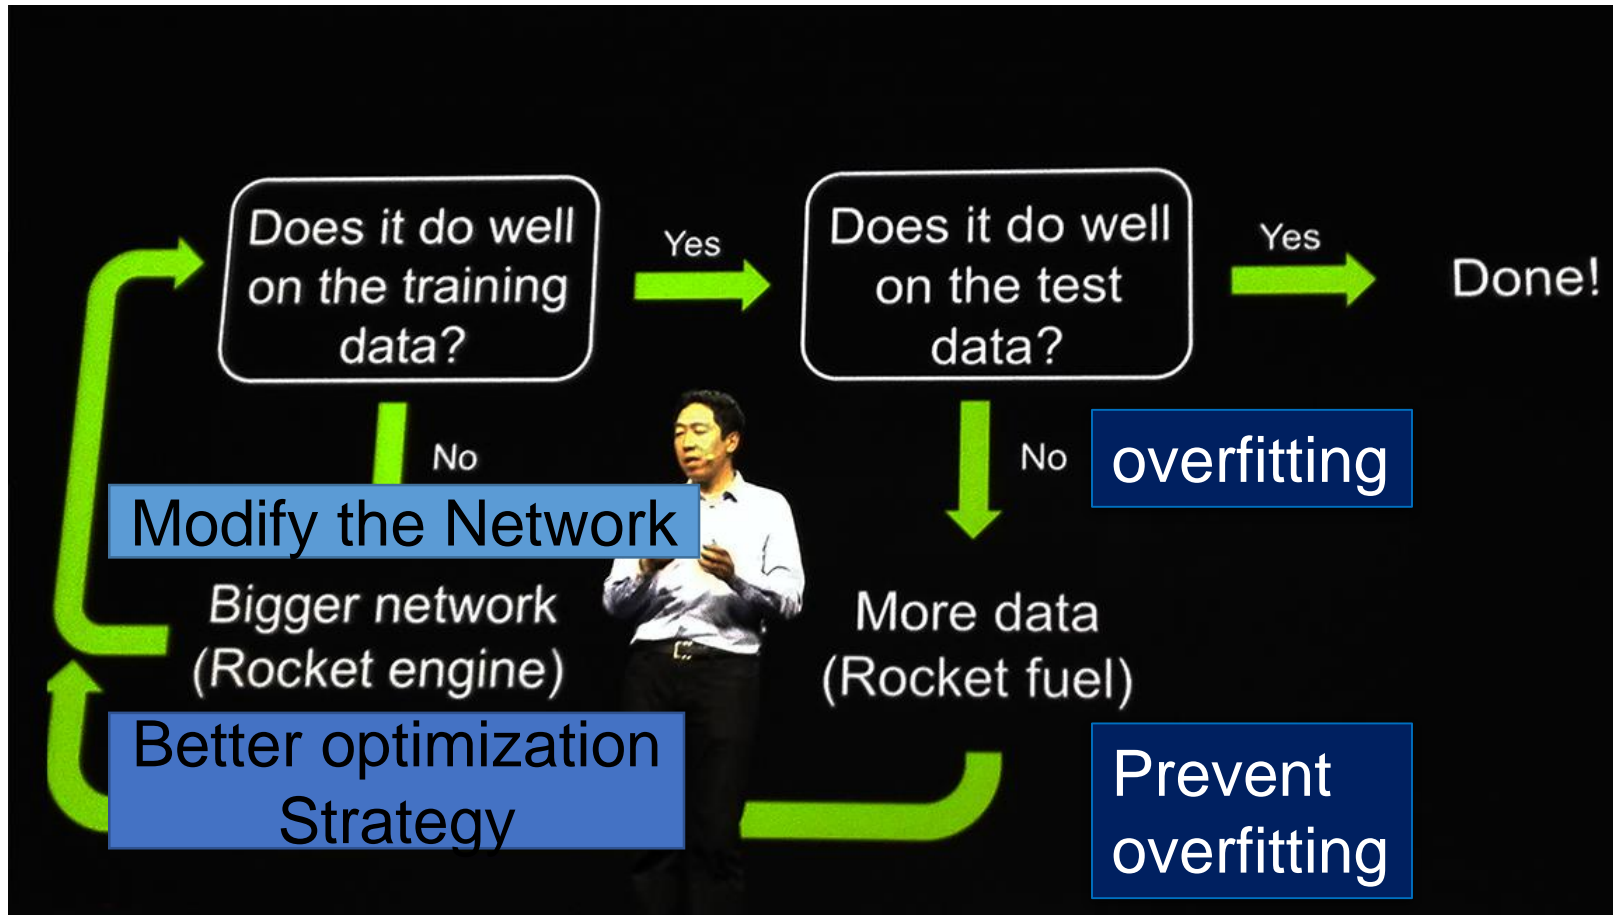


Summary: the overfitting may be due to  
**too many hidden nodes**



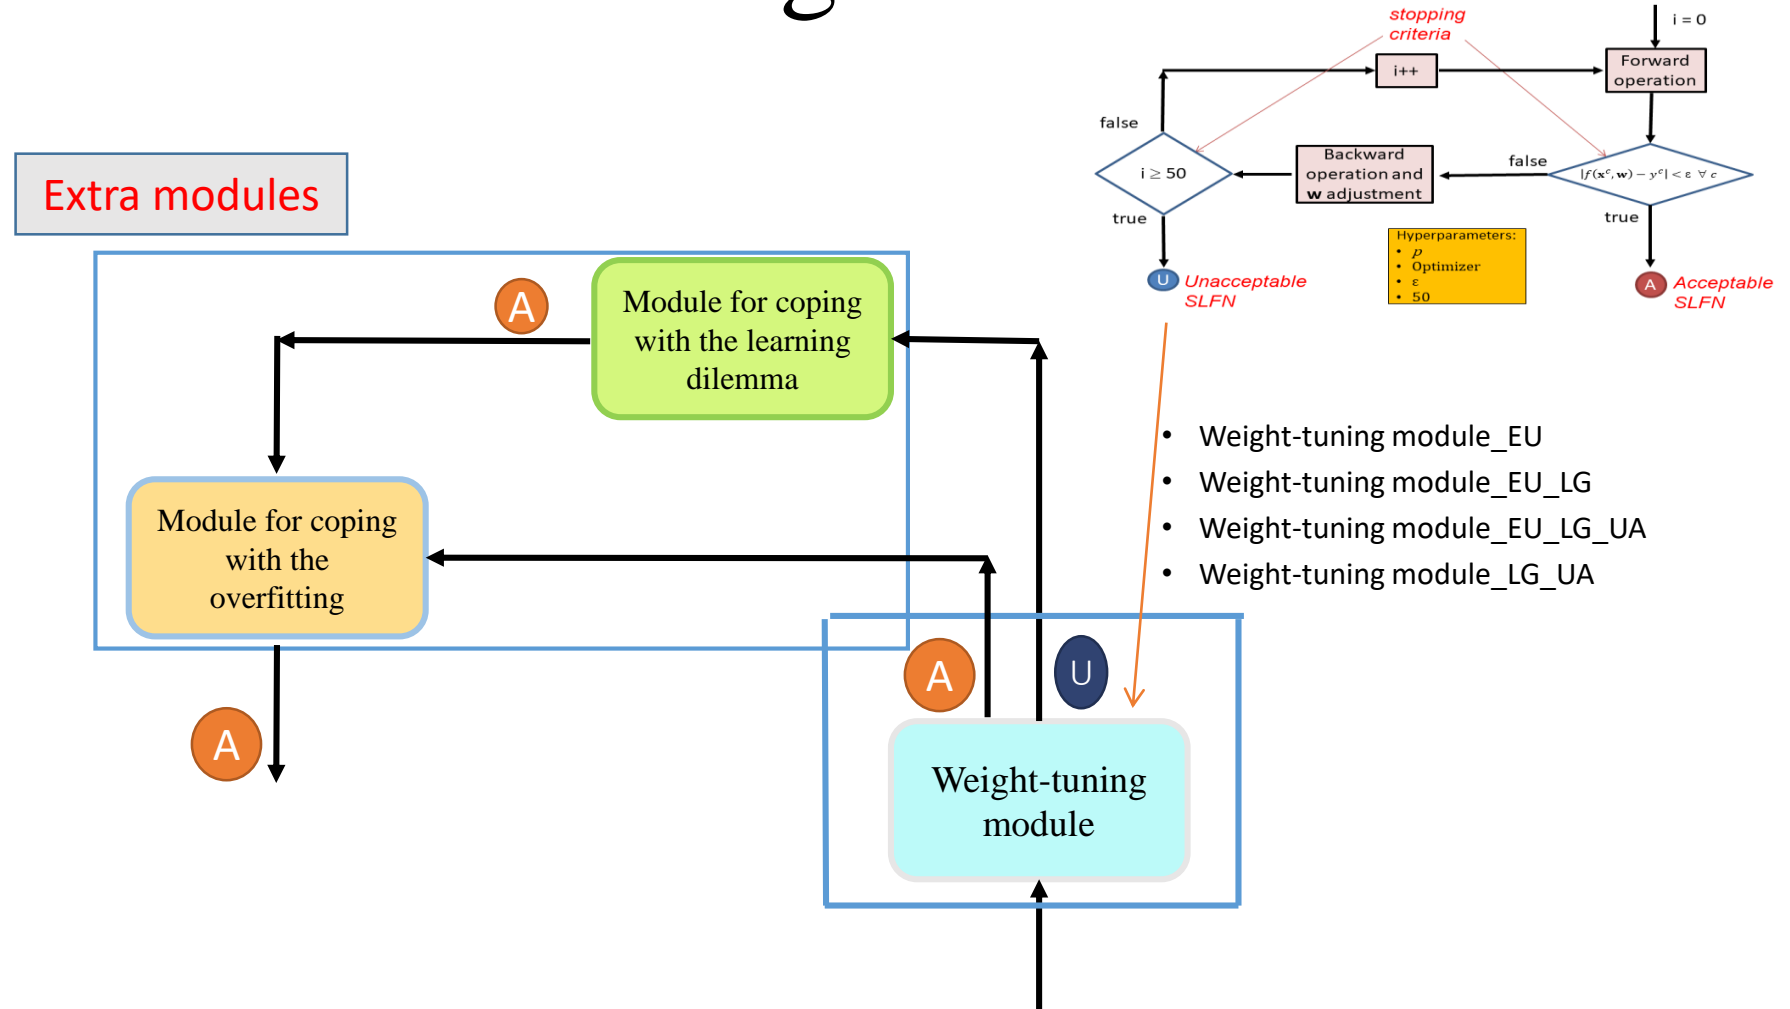
[https://www.neuraldesigner.com/images/learning/selection\\_error.svg](https://www.neuraldesigner.com/images/learning/selection_error.svg)

# Recipe for Deep Learning




<http://www.gizmodo.com.au/2015/04/the-basic-recipe-for-machine-learning-explained-in-a-single-powerpoint-slide/>

# Inferencing Issues



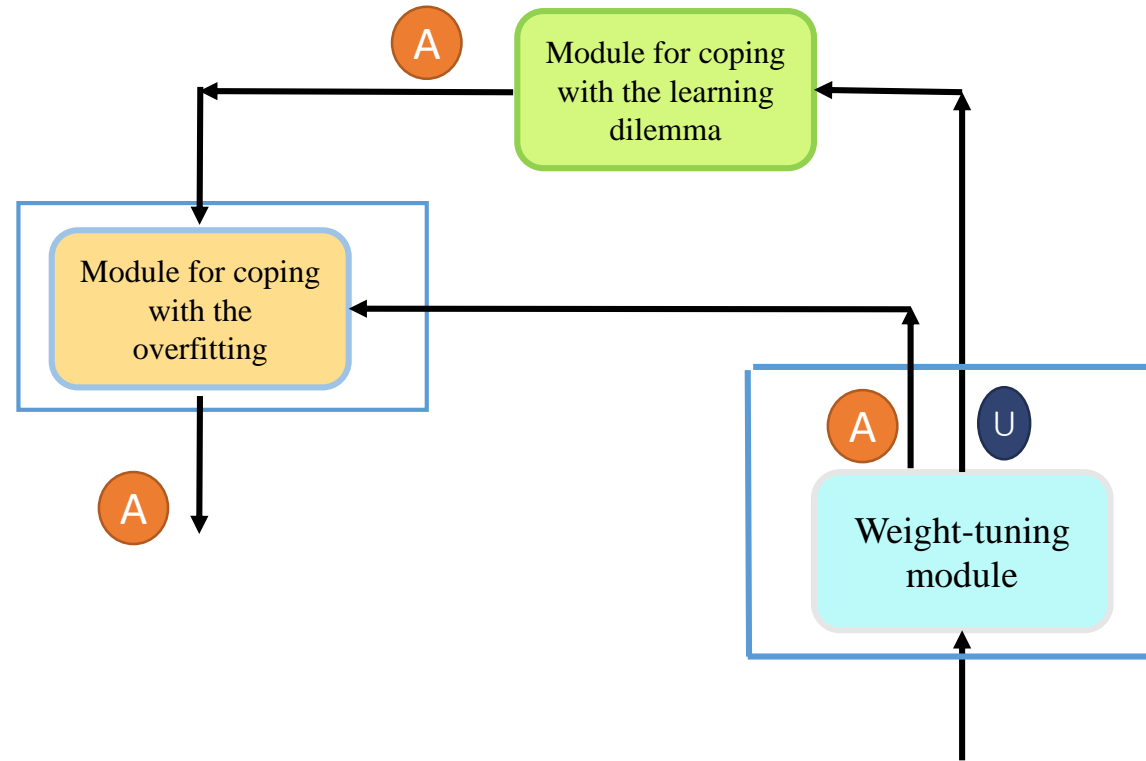
Where we are now...

# Developing a new AI system is like playing with Lego – lots of (pre-built or self-built) modules



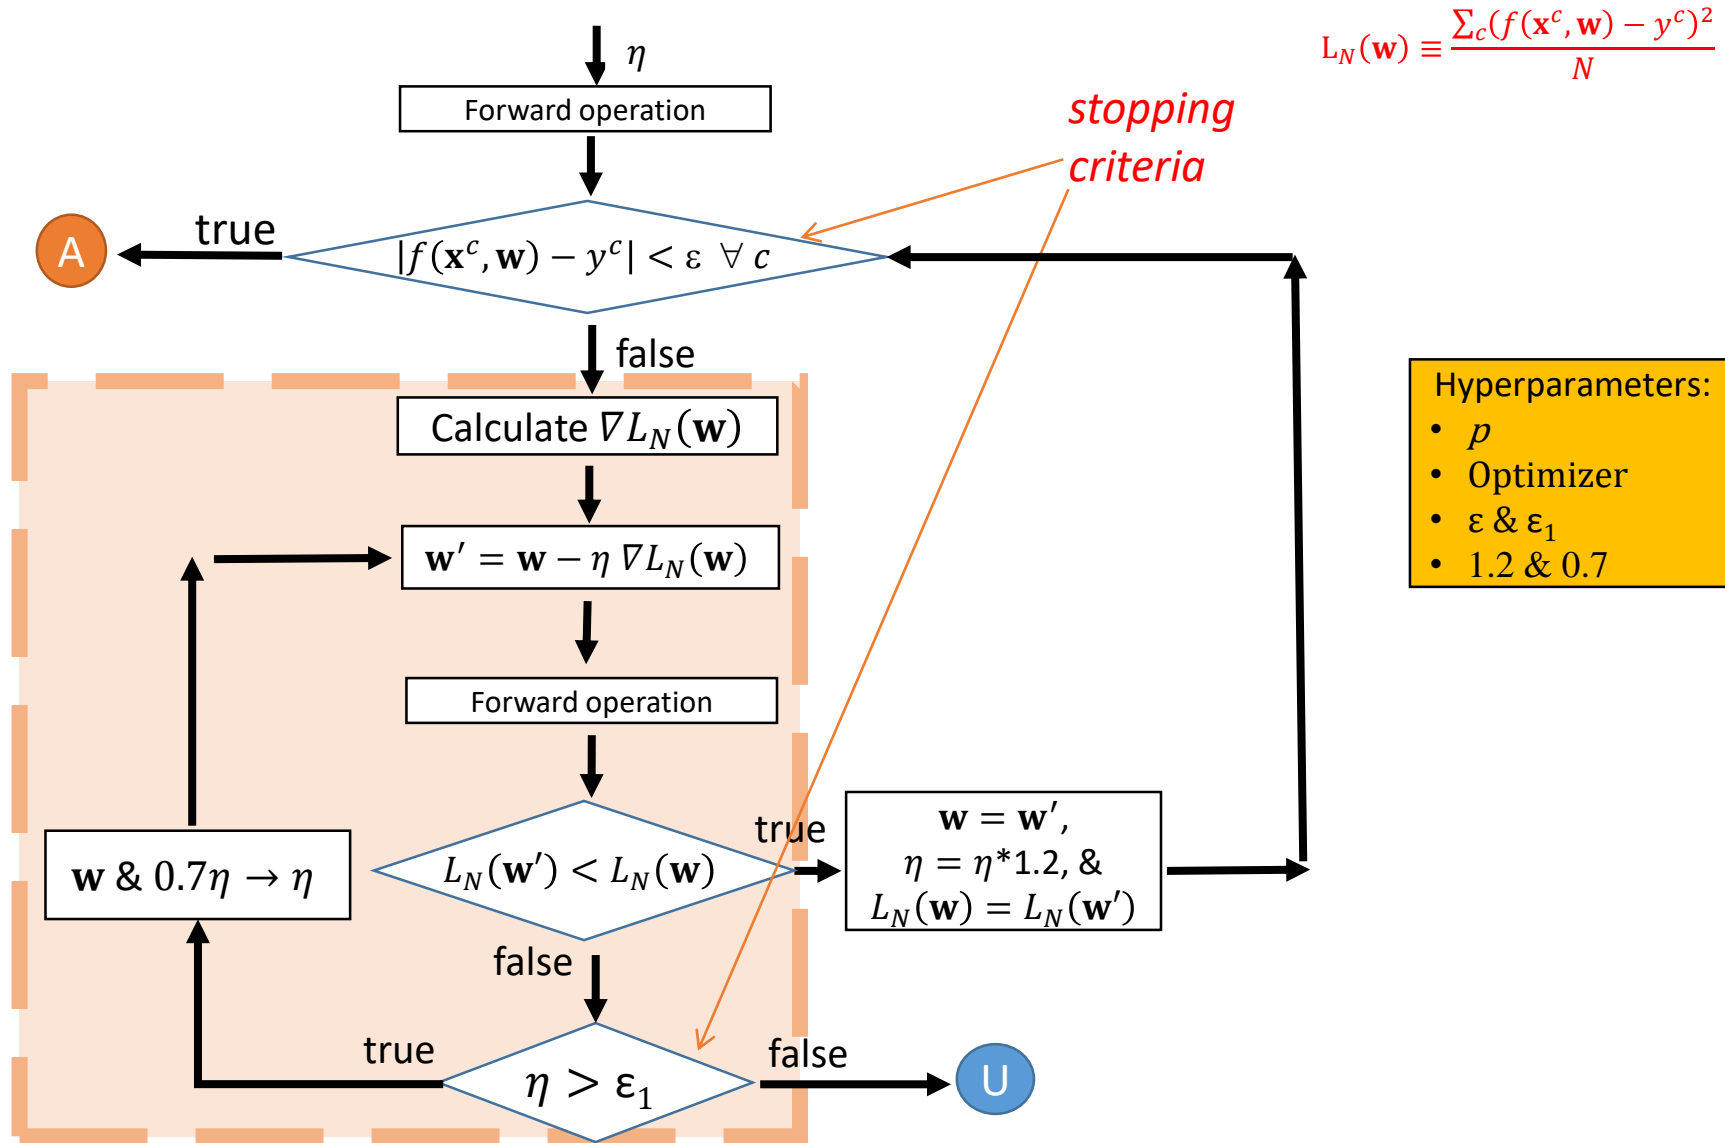
# Dealing with the overfitting due to **big weights** – the **regularizing** module

$$L_N(\mathbf{w}) \equiv \frac{\sum_c (f(\mathbf{x}^c, \mathbf{w}) - y^c)^2}{N} + \lambda \left( \sum_{i=0}^p (w_i^o)^2 + \sum_{i=1}^p \sum_{j=0}^m (w_{ij}^H)^2 \right)$$



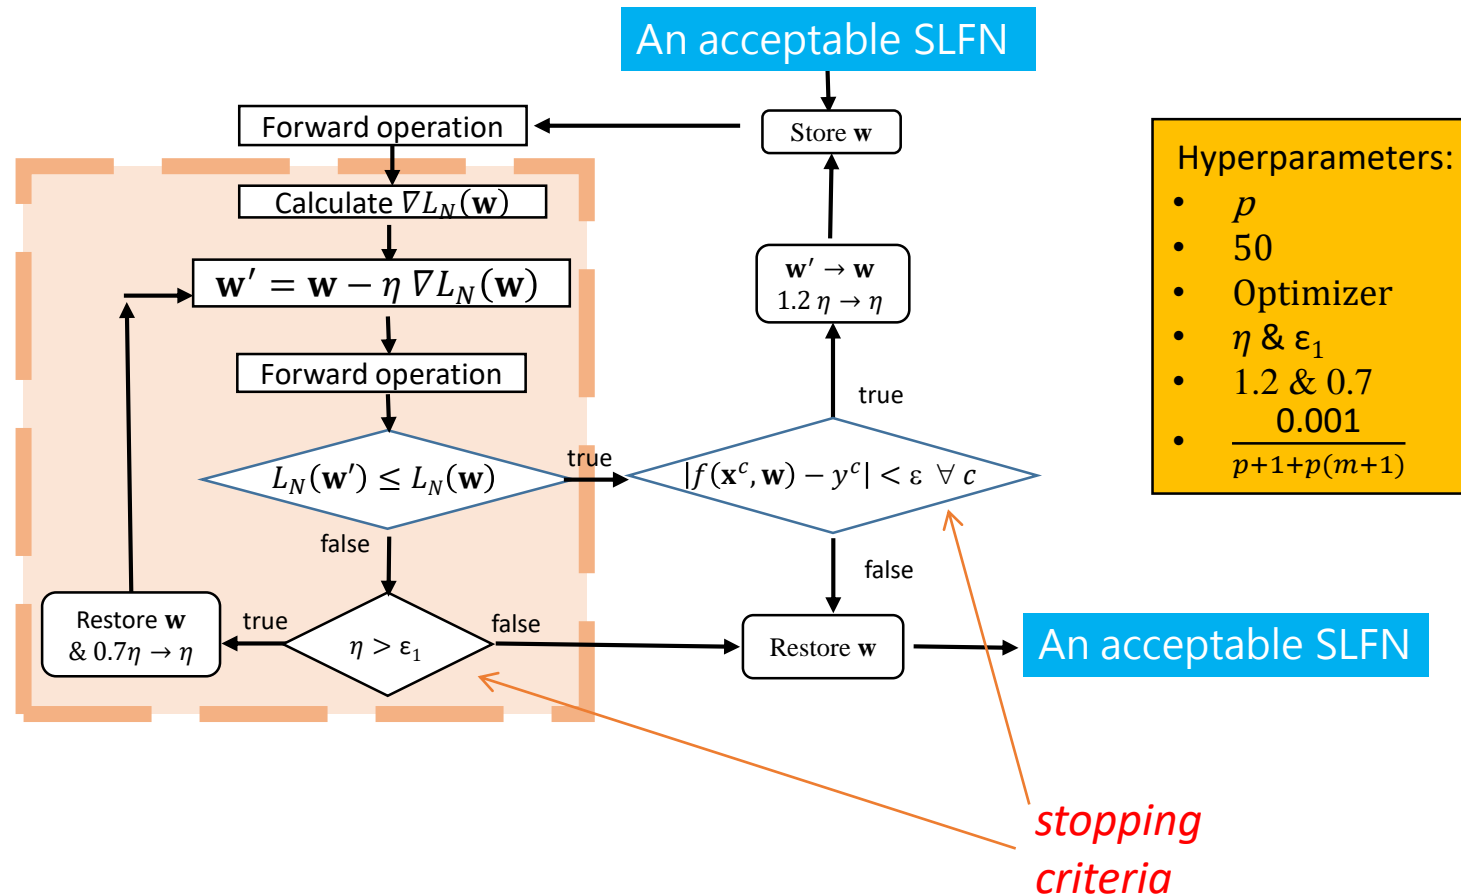
Where we are now...

# The flowchart of **weight-tuning** module\_LG\_UA

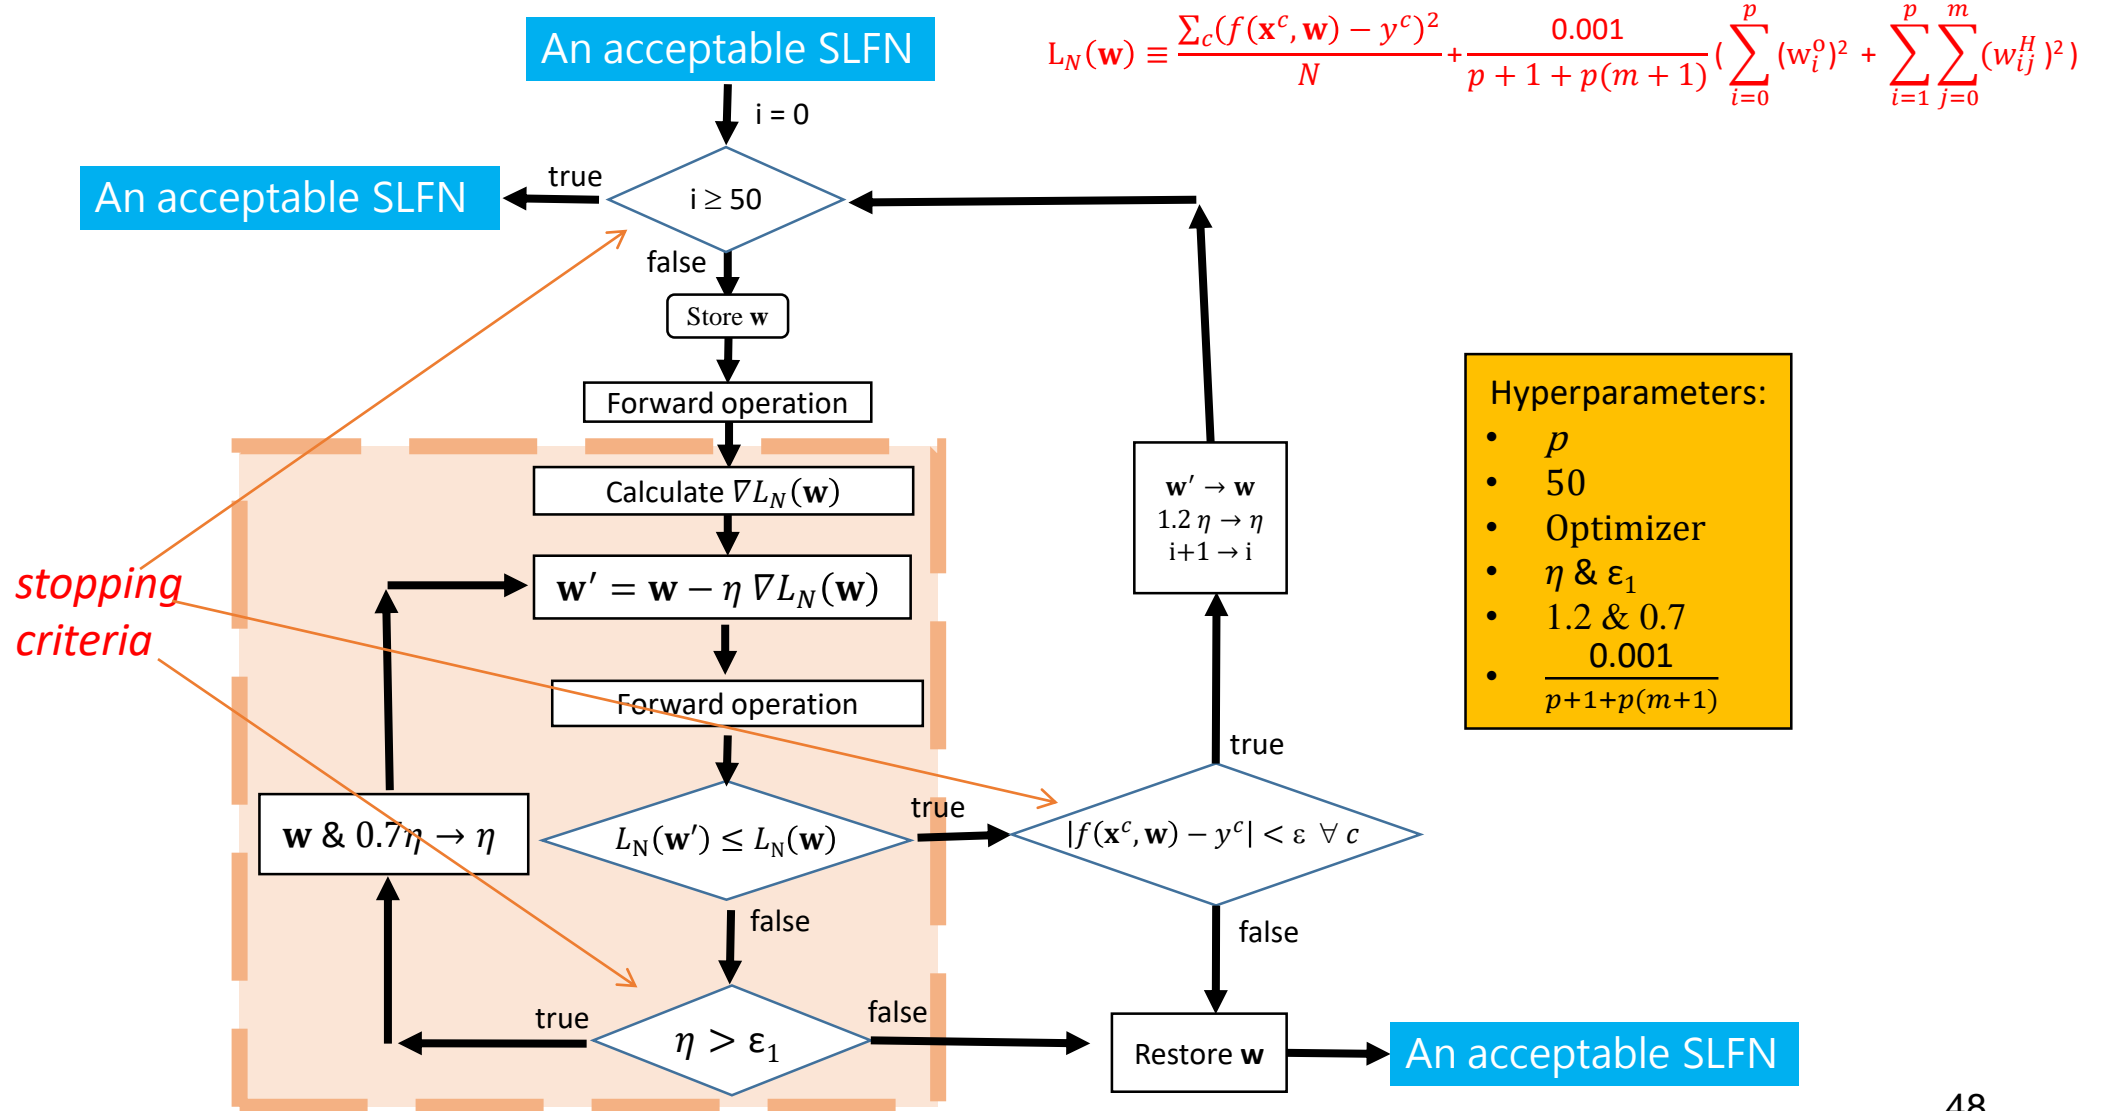


The flowchart of **regularizing** module\_LG\_UA that tries to further regularize weights of an acceptable SLFN

$$L_N(\mathbf{w}) \equiv \frac{\sum_c (f(\mathbf{x}^c, \mathbf{w}) - y^c)^2}{N} + \frac{0.001}{p+1+p(m+1)} \left( \sum_{i=0}^p (w_i^0)^2 + \sum_{i=1}^p \sum_{j=0}^m (w_{ij}^H)^2 \right)$$



# The flowchart of **regularizing** module\_EU\_LG\_UA





# The regularizing module

- The **weight-tuning** module helps tune up the weights to decrease the data error to obtain an acceptable SLFN.
- After obtaining an acceptable SLFN, the **regularizing** module helps further regularize weights of the acceptable SLFN while keeping the learning goal satisfied.
- A **well-regularized** SLFN can alleviate the **overfitting tendency**.

Q: Which optimizer does better in the regularizing module?

A: Need to conduct experiments to get the better optimizer.

# The overfitting due to **big weights**


- Adopt a regularization term in the loss function to penalize big weights:

$$L_N(\mathbf{w}) \equiv \frac{1}{N} \sum_{c=1}^N (f(\mathbf{x}^c, \mathbf{w}) - y^c)^2 + \lambda \|\mathbf{w}\|^2$$

- Decay coefficient: tiny  $\lambda$
- Regularization strength: arbitrary  $\lambda$

- The regularization strength (RS)  $\lambda$  determines how dominant the regularization is during gradient computation: Bigger  $\lambda \rightarrow$  bigger penalty for big weights
- Maybe there should be a RS scheduling like the LR scheduling. The RS should be enlarged from a tiny value.

## Learning Rate Decay



Vaswani et al, "Attention is all you need", NIPS 2017

**Step:** Reduce learning rate at a few fixed points. E.g. for ResNets, multiply LR by 0.1 after epochs 30, 60, and 90.

**Cosine:**  $\alpha_t = \frac{1}{2} \alpha_0 (1 + \cos(t\pi/T))$

**Linear:**  $\alpha_t = \alpha_0 (1 - t/T)$

**Inverse sqrt:**  $\alpha_t = \alpha_0 / \sqrt{t}$

$\alpha_0$  : Initial learning rate  
 $\alpha_t$  : Learning rate at epoch  $t$   
 $T$  : Total number of epochs

# Performance differences amongst regularizing modules

- There are two regularizing modules
  - ✓ the regularizing module\_EU\_LG\_UA
  - ✓ the regularizing module\_LG\_UA
  - ✓ the regularizing module\_EU
- What are the differences amongst these regularizing modules?

# Performance differences amongst regularizing modules

- There are two regularizing modules
  - ✓ the regularizing module\_EU\_LG\_UA  
The regularizing time length is expected
  - ✓ the regularizing module\_LG\_UA  
The regularizing time length may be much longer
  - ✓ the regularizing module\_EU  
The simplest and the regularizing time length is expected

# Regularization - In practice

**Training:** Add random noise

**Testing:** Marginalize over the noise

## Examples:

Dropout

Batch Normalization

Data Augmentation

DropConnect

Fractional Max Pooling

Stochastic Depth

Cutout / Random Crop

Mixup

- Consider dropout for large fully-connected layers
- Batch normalization and data augmentation almost always a good idea
- Try cutout and mixup especially for small classification datasets

# The regularizing module\_D0

More common: “Inverted dropout”

```
p = 0.5 # probability of keeping a unit active. higher = less dropout

def train_step(X):
    # forward pass for example 3-layer neural network
    H1 = np.maximum(0, np.dot(W1, X) + b1)
    U1 = (np.random.rand(*H1.shape) < p) / p # first dropout mask. Notice /p!
    H1 *= U1 # drop!
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
    U2 = (np.random.rand(*H2.shape) < p) / p # second dropout mask. Notice /p!
    H2 *= U2 # drop!
    out = np.dot(W3, H2) + b3

    # backward pass: compute gradients... (not shown)
    # perform parameter update... (not shown)


def predict(X):
    # ensembled forward pass
    H1 = np.maximum(0, np.dot(W1, X) + b1) # no scaling necessary
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
    out = np.dot(W3, H2) + b3
```

test time is unchanged!

# The regularizing module\_BN

## Batch Normalization

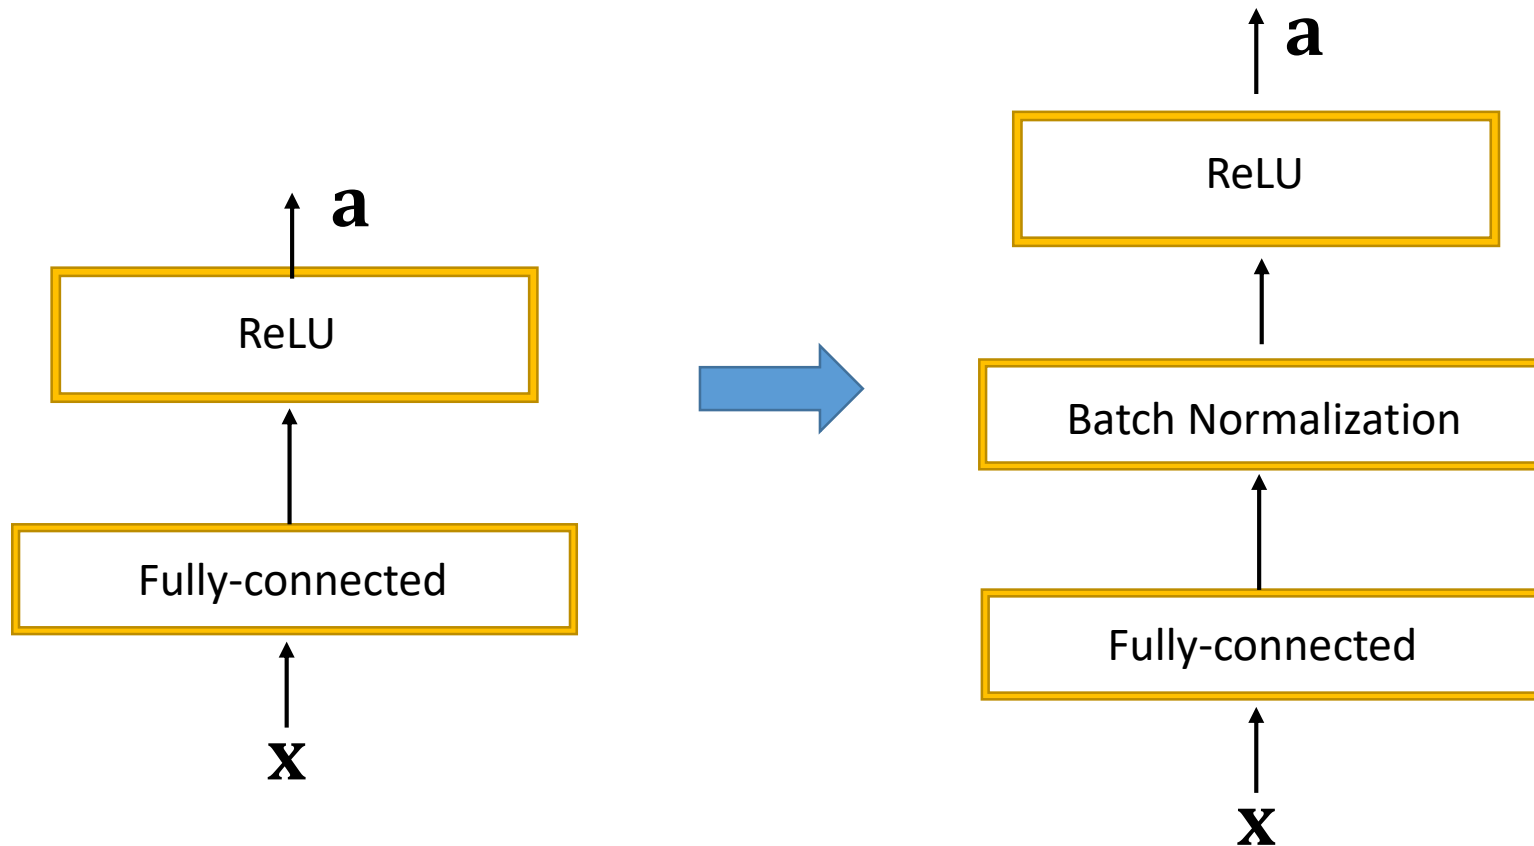
[Ioffe and Szegedy, 2015]



- Makes deep networks **much** easier to train!
- Improves gradient flow
- Allows higher learning rates, faster convergence
- Networks become more robust to initialization
- Acts as regularization during training
- Zero overhead at test-time: can be fused with conv!
- Behaves differently during training and testing: this is a very common source of bugs!

# The regularizing module\_BN

In the regularizing module\_BN, the batch normalization operation is inserted after the FC layer and before the nonlinearity layer.






# Homework 3\_1

- Rewrite the code of **weight-tuning module\_LG\_UA** stated in page 46 to make the coding of **regularizing module\_LG\_UA** stated in page 47.
- Rewrite the code of **regularizing module\_LG\_UA** stated in page 47 to make the coding of **regularizing module\_EU\_LG\_UA** stated in page 48.
- Make the coding of **regularizing module\_DO** stated in page 54.
- Make the coding of **regularizing module\_BN** stated in page 56.

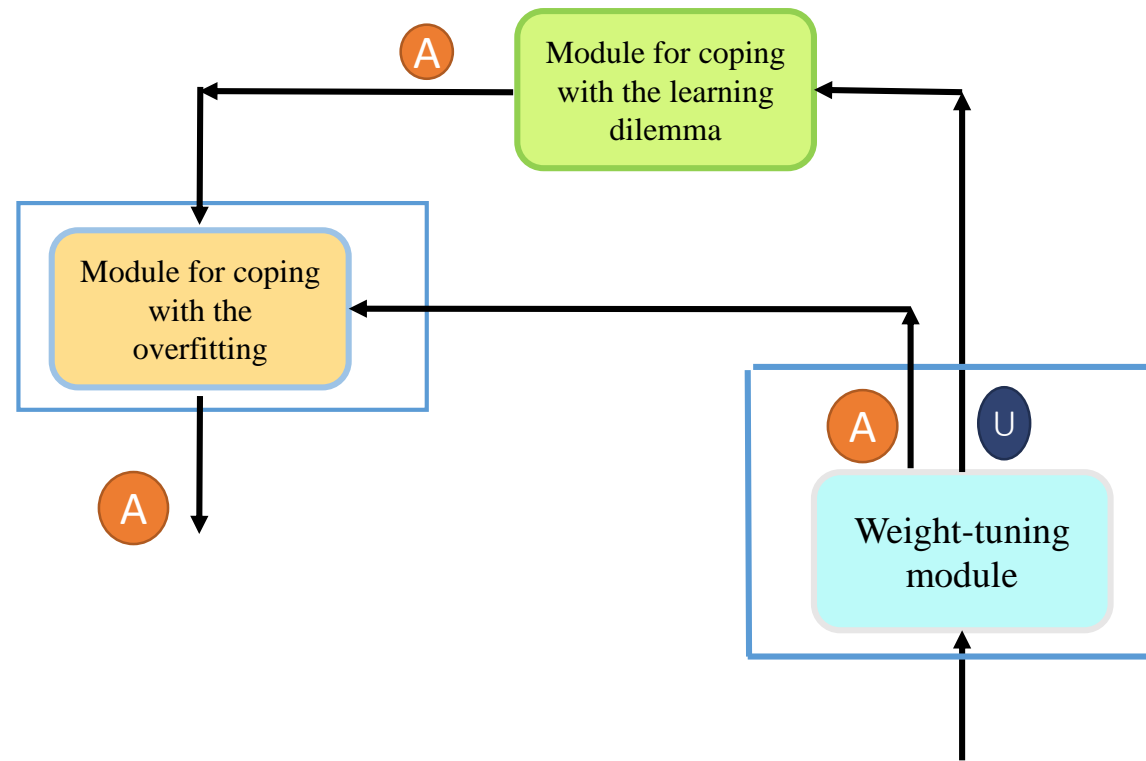
Where we are now...

# Overfitting due to **too many hidden nodes**



[https://www.neuraldesigner.com/images/learning/selection\\_error.svg](https://www.neuraldesigner.com/images/learning/selection_error.svg)

# Dealing with the overfitting due to **big weights** and **too many hidden nodes** – the **reorganizing** module



# Regularization

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Data loss}} + \underbrace{\lambda R(W)}_{\text{Regularization}}$$

**Data loss:** Model predictions should match training data


**Regularization:** Prevent the model from doing *too* well on training data

**Occam's Razor:** Among multiple competing hypotheses, the simplest is the best,  
William of Ockham 1285-1347


# irrelevant hidden nodes & potentially irrelevant hidden nodes

- Develop the **reorganizing module** that helps regularize weights of an acceptable SLFN while keeping the learning goal satisfied as well as identify and remove the *potentially irrelevant hidden node*.
- The hidden node that can be pruned without making the learning goal unsatisfied is an *irrelevant hidden node*. (Tsaih, 1993)
- For the SLFN with the  $\mathbf{w}$ , the  $k^{\text{th}}$  hidden node is *potentially irrelevant* if the learning goal can be accomplished via minimizing  $L_N(\mathbf{w}'_k)$ , where  $\mathbf{w}'_k \equiv \mathbf{w} - \{w_k^O, w_{k0}^H, \mathbf{w}_k^H\}$  and  $f(\mathbf{x}^c, \mathbf{w}'_k) \equiv w_0^O + \sum_{i \neq k} w_i^O a_i^c \quad \forall c$ . (Tsaih, 1993)

**W**



$$\mathbf{w}'_k \equiv \mathbf{w} - \{w_k^o, w_{k0}^H, \mathbf{w}_k^H\}$$



# The reorganizing module that helps regularize weights of an acceptable SLFN and examines its hidden nodes one by one

Note that there are two optimizers:

One for the regularizing purpose

Another for the pruning purpose

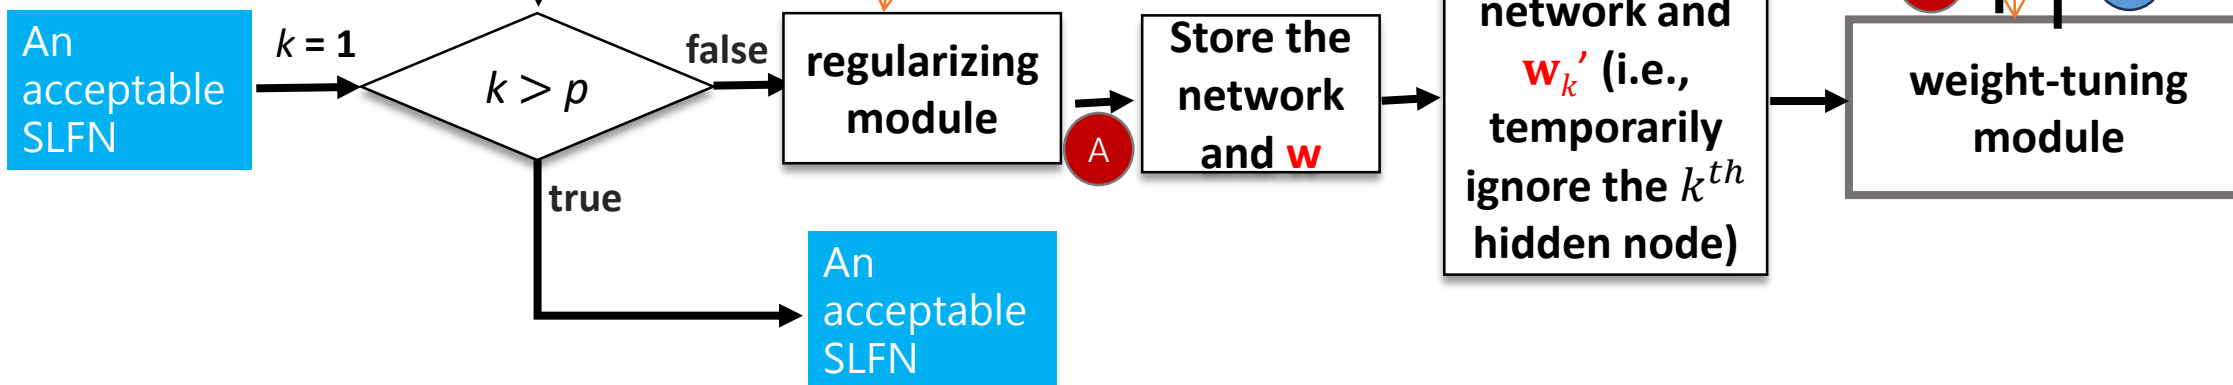
$$L_N(\mathbf{w}) \equiv \frac{\sum_c (f(\mathbf{x}^c, \mathbf{w}) - y^c)^2}{N} + \frac{0.001}{p+1+p(m+1)} \left( \sum_{i=0}^p (w_i^0)^2 + \sum_{i=1}^p \sum_{j=0}^m (w_{ij}^H)^2 \right)$$

$$L_N(\mathbf{w}) \equiv \frac{\sum_c (f(\mathbf{x}^c, \mathbf{w}) - y^c)^2}{N}$$

Will you use different optimizers?

Hyperparameters:

- The epoch constraint
- Two optimizers
- Two  $\eta$
- Two sets of  $\varepsilon$  &  $\varepsilon_1$
- 1.2 & 0.7
- $\frac{0.001}{p+1+p(m+1)}$





# The reorganizing module **\_ALL\_r\_EU\_LG\_UA\_w\_EU\_LG\_UA**

Note that there are two optimizers:

One for the regularizing purpose

Another for the pruning purpose

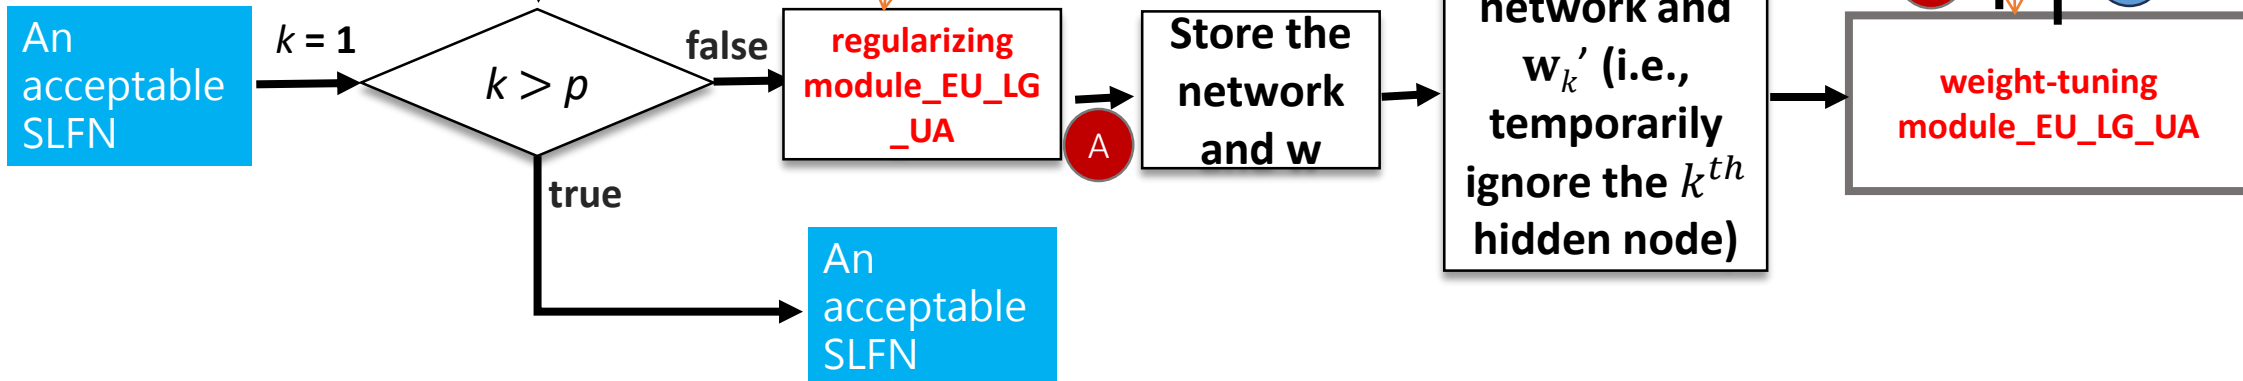
$$L_N(\mathbf{w}) \equiv \frac{\sum_c (f(\mathbf{x}^c, \mathbf{w}) - y^c)^2}{N} + \frac{0.001}{p+1+p(m+1)} \left( \sum_{i=0}^p (w_i^0)^2 + \sum_{i=1}^p \sum_{j=0}^m (w_{ij}^H)^2 \right)$$

$$L_N(\mathbf{w}) \equiv \frac{\sum_c (f(\mathbf{x}^c, \mathbf{w}) - y^c)^2}{N}$$

Will you use different optimizers?

Hyperparameters:

- The epoch constraint
- Two optimizers
- Two  $\eta$
- Two sets of  $\varepsilon$  &  $\varepsilon_1$
- 1.2 & 0.7
- $\frac{0.001}{p+1+p(m+1)}$



# More ideas for the reorganizing module that examines merely some hidden nodes one by one

1. The reorganizing module **\_R3\_r\_EU\_LG\_UA\_w\_EU\_LG\_UA** that randomly **picks up 3** hidden nodes and examines whether they are potentially irrelevant. Remove potentially irrelevant hidden nodes identified within the process.
2. The reorganizing module **\_MAW\_r\_EU\_LG\_UA\_w\_EU\_LG\_UA** that uses  $k = \arg \min_i |w_i^o|$  to pick up a hidden node and examines whether it is potentially irrelevant. If yes, remove it and then repeat the process; otherwise, stop the process.
3. The reorganizing module **\_PCA\_r\_EU\_LG\_UA\_w\_EU\_LG\_UA** that uses **PCA** to pick up a hidden node and examines whether it is potentially irrelevant. If yes, remove it and then repeat the process; otherwise, stop the process.
4. The reorganizing module **\_ETP\_r\_EU\_LG\_UA\_w\_EU\_LG\_UA** that calculates the **entropy** of each hidden node and then, based on the obtained entropy, picks up a hidden node and examines whether it is potentially irrelevant. If yes, remove it and then repeat the process; otherwise, stop the process.

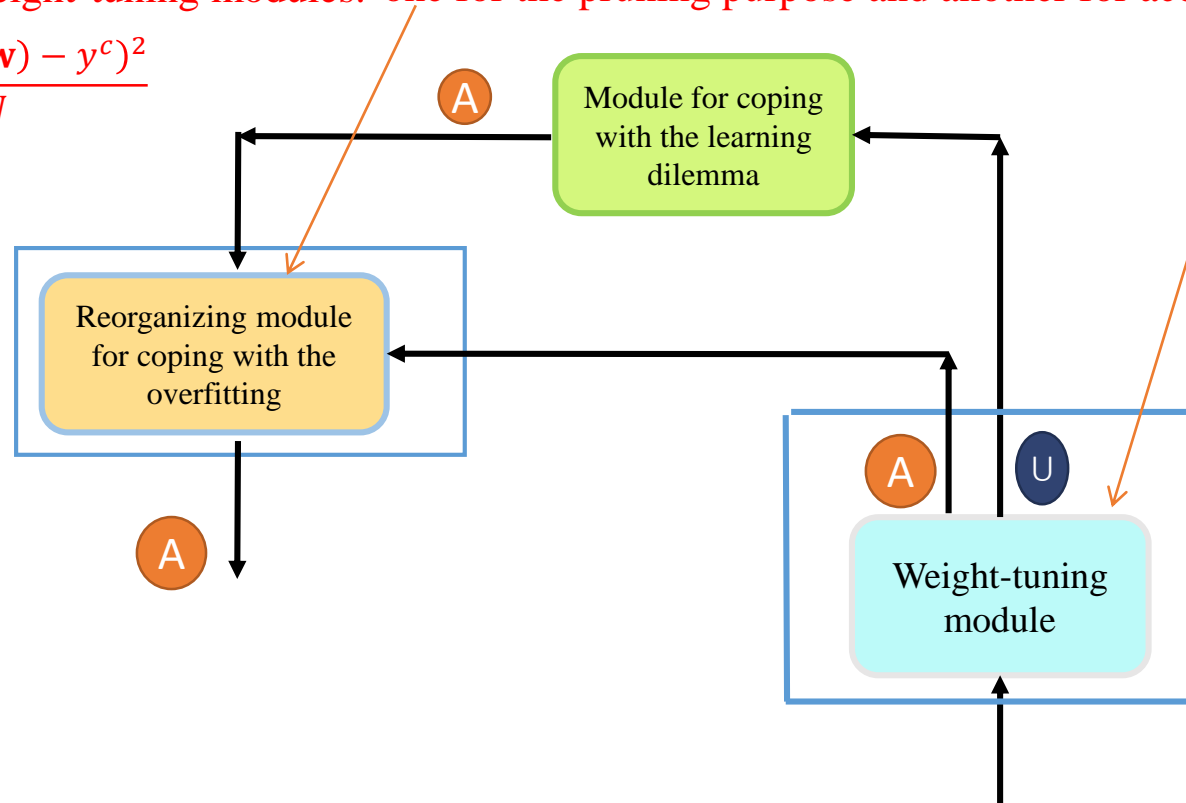
In information theory, the **entropy** of a random variable is the average level of "information", "surprise", or "uncertainty" inherent in the variable's possible outcomes. Given a discrete random variable  $\mathbf{X}$ , with possible outcomes  $x_1, \dots, x_n$ , which occur with probability  $P(x_1), \dots, P(x_n)$ , the entropy of  $\mathbf{X}$  is formally defined as:  $H(\mathbf{X}) = - \sum_{i=1}^n P(x_i) \log(P(x_i))$ . ([Entropy - Wikipedia](#))

5. Your idea?

# The **weight-tuning** module and the **reorganizing** module

Note that there are two weight-tuning modules: one for the pruning purpose and another for accomplishing the learning goal.

$$L_N(\mathbf{w}) \equiv \frac{\sum_c (f(\mathbf{x}^c, \mathbf{w}) - y^c)^2}{N}$$



Will you use different weight-tuning modules for the pruning purpose and for accomplishing the learning goal?

## Homework 3

- Make the coding of AI system stated in page 67 without the module for coping with the learning dilemma.
- Use the reorganizing module `_ALL_r_EU_LG_UA_w_EU_LG_UA` stated in page 65.
- You may pick up one of the following weight-tuning modules:
  - ✓ the weight-tuning module `_EU`
  - ✓ the weight-tuning module `_EU_LG`
  - ✓ the weight-tuning module `_EU_LG_UA`
  - ✓ the weight-tuning module `_LG_UA`
- Note that the learning goals used in the weight-tuning module, the regularizing module, and the reorganizing module should be the same.